

# Recurrent Neural Networks

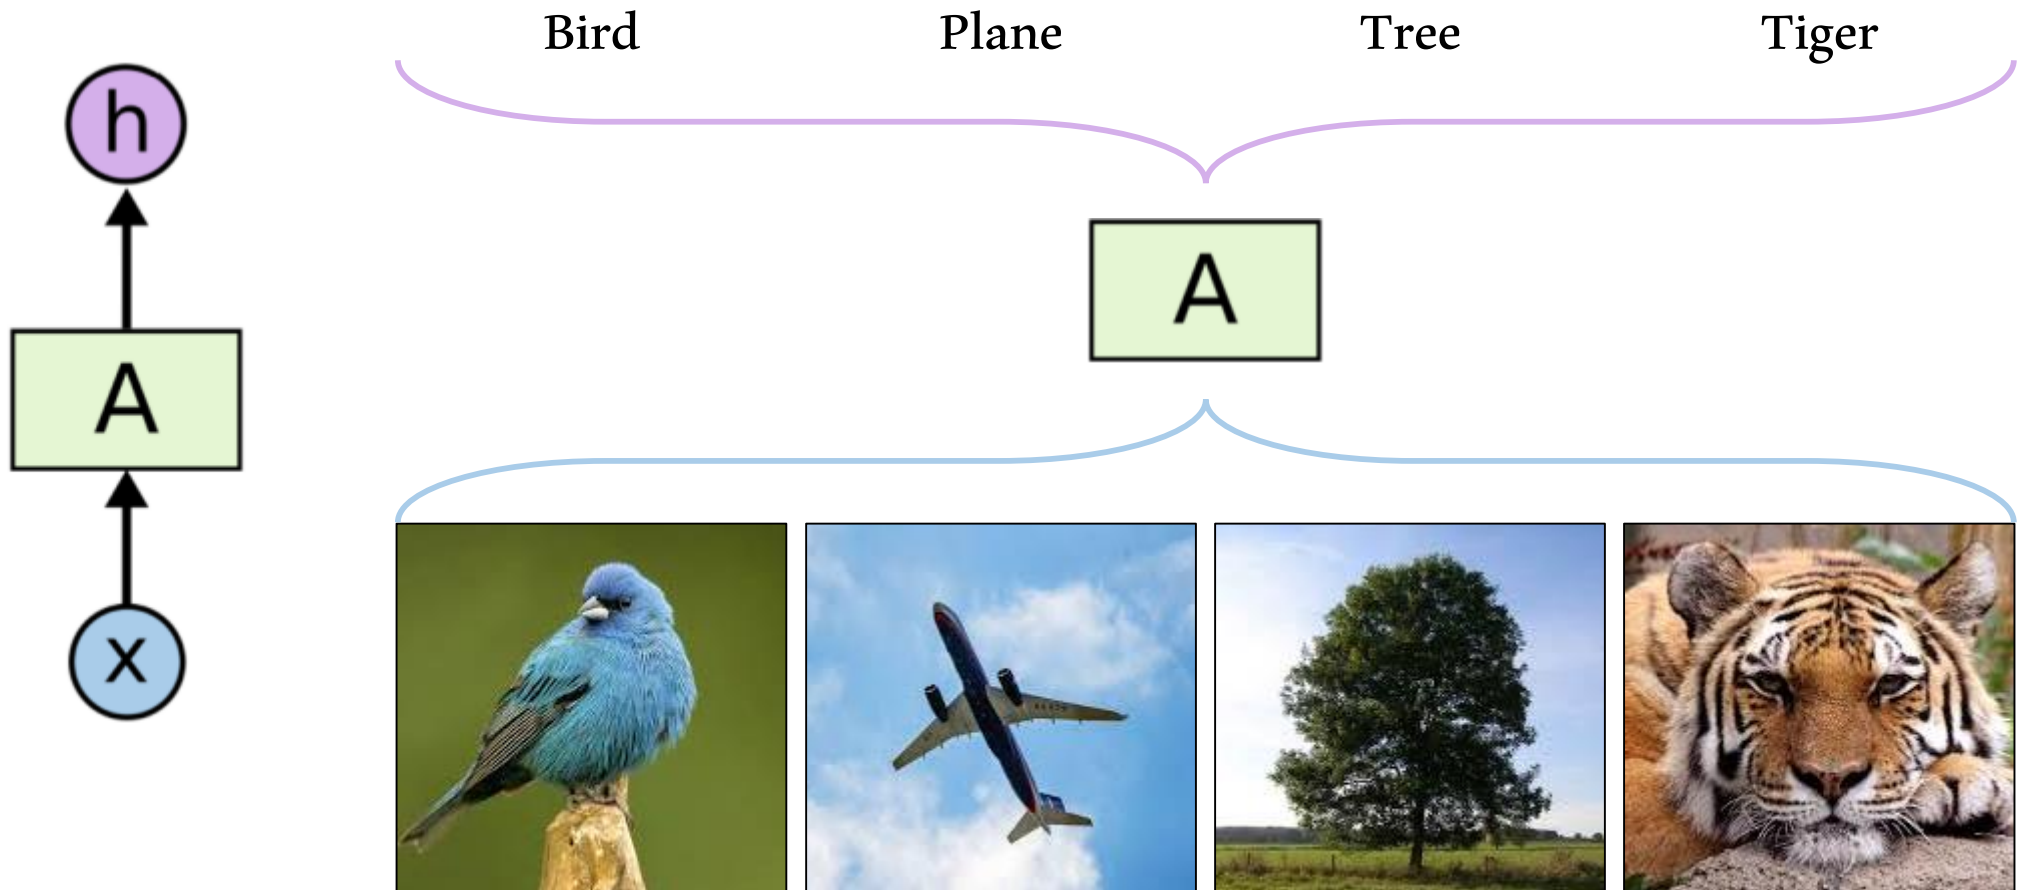
Paola Ruiz & Natalia Valderrama

Tutor: Laura Daza

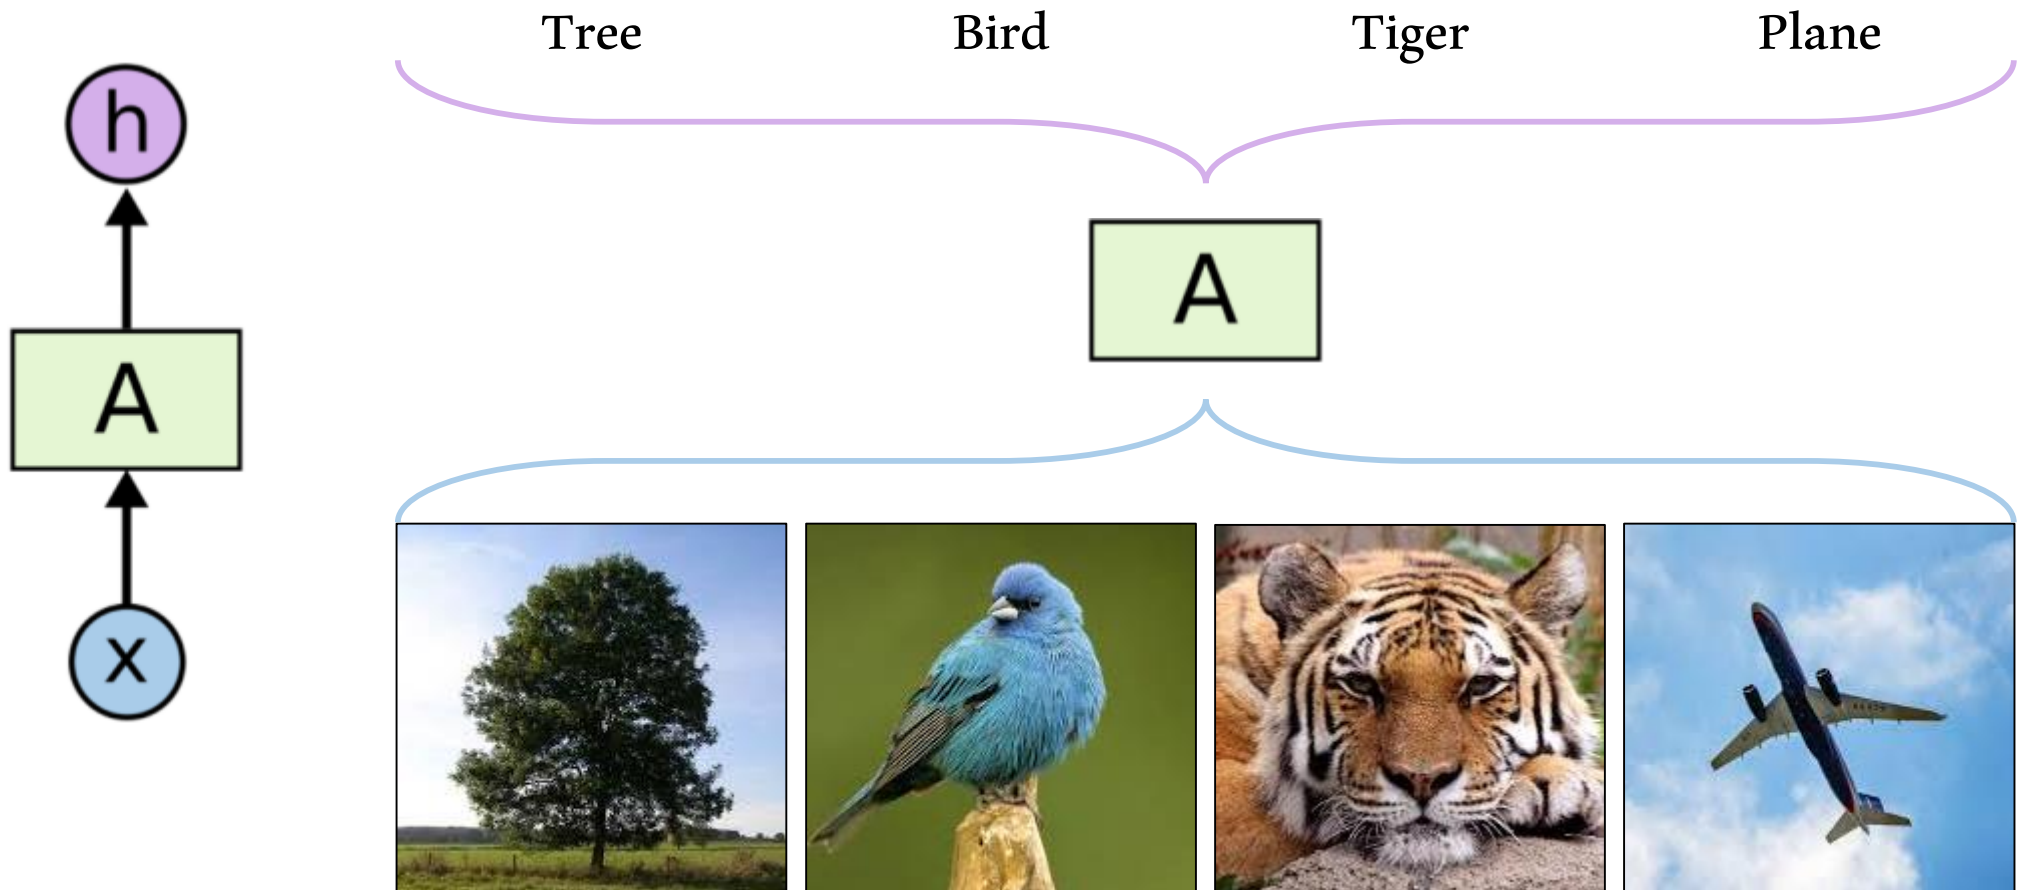
# Content List

- ✓ Introduction
- ✓ RNN structure
- ✓ Types
- ✓ Backpropagation through time
- ✓ Vanishing and exploding Gradient Problem
- ✓ LSTM
- ✓ GRU
- ✓ SRU
- ✓ Variation diagrams:
  - ✓ Multilayer
  - ✓ Bidirectional
  - ✓ Encoder-Decoder
  - ✓ Attention Layer
- ✓ Papers:
  - ✓ SMT
  - ✓ Biometrics
- ✓ Tutorial
- ✓ Homework

# Introduction



# Introduction



# Introduction

"I've got to think that that was unethical," Joshua said.

"Josh, faking demonic possession is like a mustard seed."

"How is it like a mustard seed?"

"You don't know, do you? Doesn't seem at all like a mustard seed, does it? Now you see how we all feel when you liken things unto a mustard seed? Huh?"

– Christopher Moore, *Lamb: The Gospel According to Biff, Christ's Childhood Pal*


"I also felt guilty about the three pens I'd stolen, but only for a second.

And since there was no convenient way to give them back, I stole a bottle of ink before I left."

– Patrick Rothfuss, *The Name of the Wind*

"It wasn't even a good note. 'If you are reading this I am probably dead.' What sort of a note is that?"

– Patrick Rothfuss, *The Name of the Wind*



Protein PFF0165c


MSNKKRSKNENDESTSLPENSELLIEYIHNLSCLNVYRREIQEKNKYISIIKNDLSFHEC  
ILTNVNVVSVFNNDDLNLCCNNEQKEEGEEI IKQRNIGDEINEYNNLTKLQNDENIKNNNM  
IKEDLEDDANQNILMKSPYYNIENFLQVFLKYINKKKKKVKVKVDEGKKEKIEDKKYEQDD  
EENEEEEEEEEEEENKEDEEFFKTFVSFNLYHNNNEKNISYDKNLVQENDNKDEAR  
GNDNMCNYDIHNERGEMLDKGKSYSGDEKINTSDNAKSCSGDEKVTSDNGKSYDYVKNES  
EEQEEKENMLNNKRSLECNPNKAKKICFSLEEKIGTVQSVKLKEYNELSKENIEKNKHDDN  
NICNYLSHNEGENVIEREDKLFNKLNNKNYRNEEEKKNQINFDYLKKIKNNQDVFEETIQ  
KCFLINLKKTLNLINKIMYLKNVEFRKYNLDYIRKINYEKCFYKKNYIDIKKISELQKDNE  
SLKIQVDRLEKKKATLIYKLNNNDNIRKHILDNNIKDYQNGIDNSKVSFYFDEGENPYNNRNN  
YRTDNKNSDDNNNNNNYNNYNSDDNYSNEDNEYNNGNYRFRNNYKDDSLNEDDVKKNPLK  
VCHKINSDSNIFVNFENIITKQNIHSEPPFRNLLKESNELYITLKEKEKENIILKNEILKME  
NKKDEEYEHLLNNTIEDKKELTRSIKELEINMMTCNMEKDKISNKVNTLEYEINVLKNIDKN  
QTMQLQKQENDILKMKLYIEKLKLSEKNLKDKIILLENKDKMLSGIHKDNSPFNEESKSEE  
GKIQLRDIQNDNEKYDEKKRFKELFIENQKLKEELNKKRNVEEELHSLRKNYNIINEEIE  
EITKEFEKKQEQVDEMILQIKNKELELLDKFNNKMKNAYVEEKLKELKNTYEEKMKHNNIY  
KKHDDFVNIYLNLFQARKNAILSDSQREEQMNLFIKLKDYDIIFQKKIELTDILKNVYDC  
NKKLIGHCQDLEKENSTLQNKLSNEIKNSKMLSKNLSKNSDDHLLIEENNELRRRLICSVCM  
ENFRNYIIKCGHIYCNCCIFNNLKTNRKCPQCKVPFDKDLQKIFLD

Schematic Representation




# Introduction

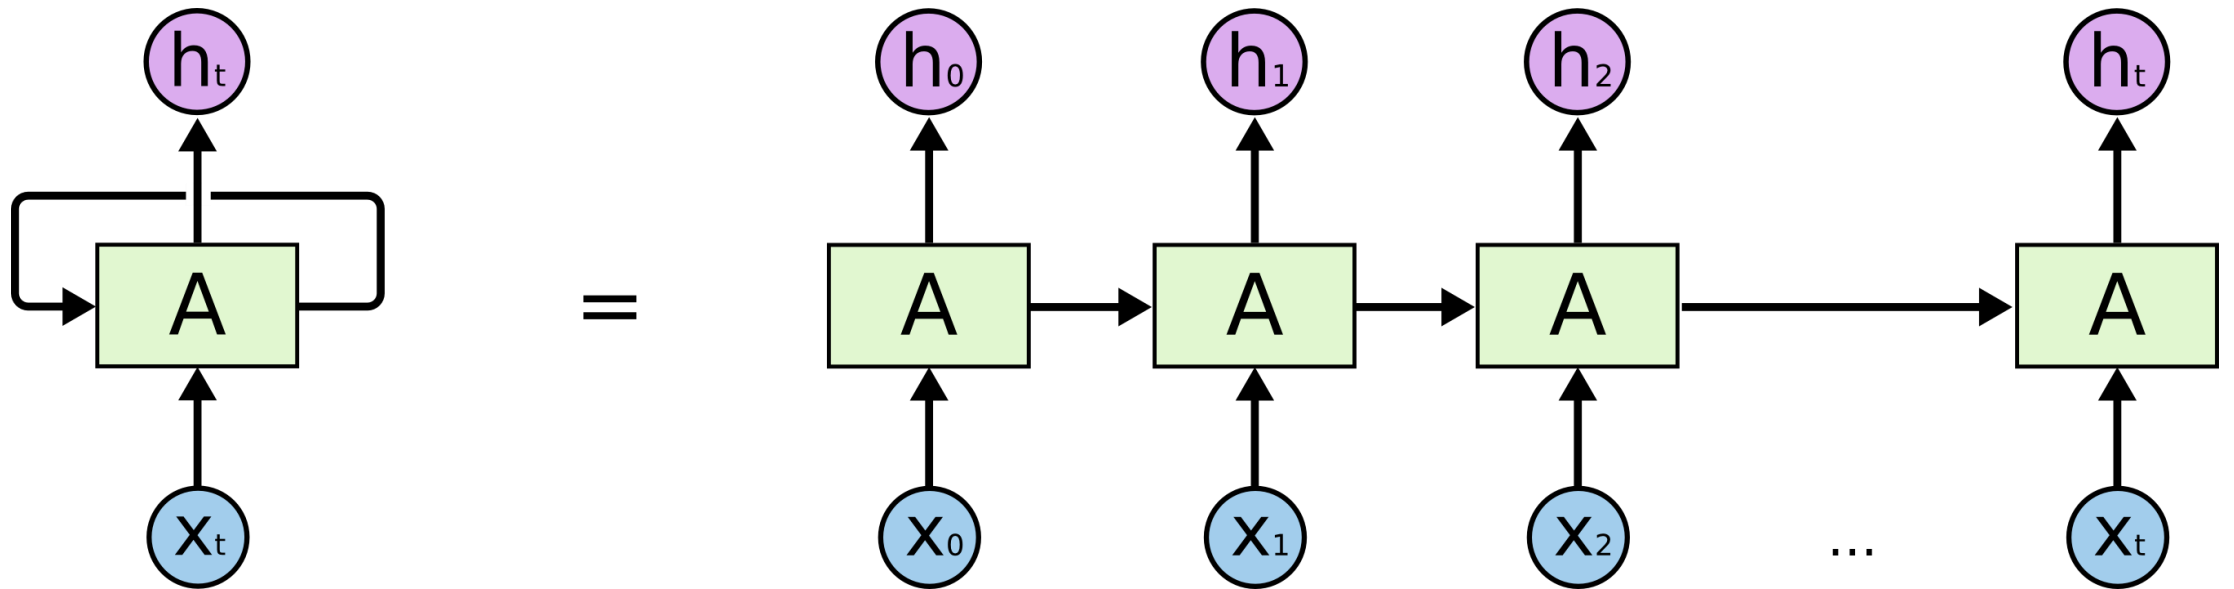
Vanilla Neural Network



Recurrent Neural Network (RNN)



# Introduction



Unrolled Recurrent Neural Network (RNN)

# Introduction

RNNs are nice and everything, but...

- How far should our “memory” go?

“I've got to think that that was unethical," Joshua said.

– Christopher Moore, *Lamb: The Gospel According to Biff, Christ's Childhood Pal*

- If the sequence is too long, wouldn't it be too heavy to store all that information?
- How do we train them?





# Recurrent Neural Networks

Paola Ruiz & Natalia Valderrama


Tutor: Laura Daza

# Recurrent Neural Networks (RNN)



Graphic  
Representation

# Recurrent Neural Networks (RNN)



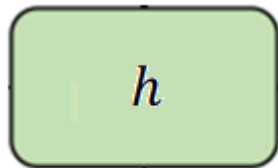
Recurrent Core  
Cell

# Recurrent Neural Networks (RNN)

$s$




$h$



- Hidden State
- Compute recurrent relation with a function  $f_w$

$$h_t = f_w(h_{t-1}, x_t)$$

# Recurrent Neural Networks (RNN)





- Input  $X$  with weights  $U$  ( $W_{xh}$ )
- Computes  $h$  with weights  $W$  ( $W_{hh}$ )

$$h_t = f_w(h_{t-1}, x_t)$$

$$h_t = \tanh(W_{hh}h_{t-1} + W_{xh}x_t)$$


# Recurrent Neural Networks (RNN)



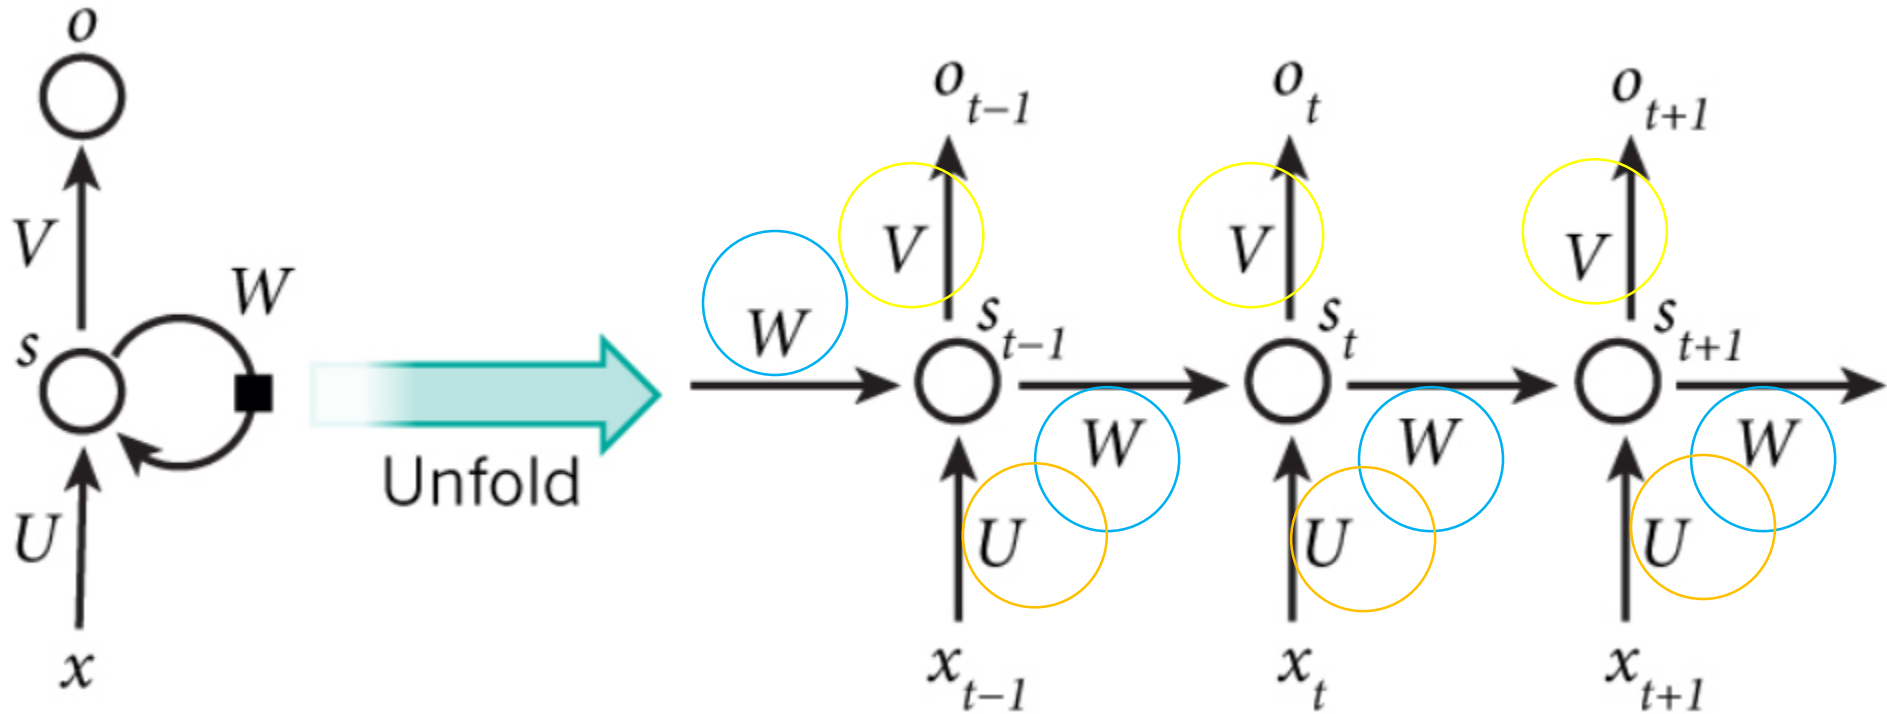
- Output  $o$  ( $\hat{y}$ ) with weights  $V$  ( $W_{hy}$ )
- Additional fully connected layers that read  $h$  to produce an output.

$$y_t = W_{hy}h_t$$

# Recurrent Neural Networks (RNN)

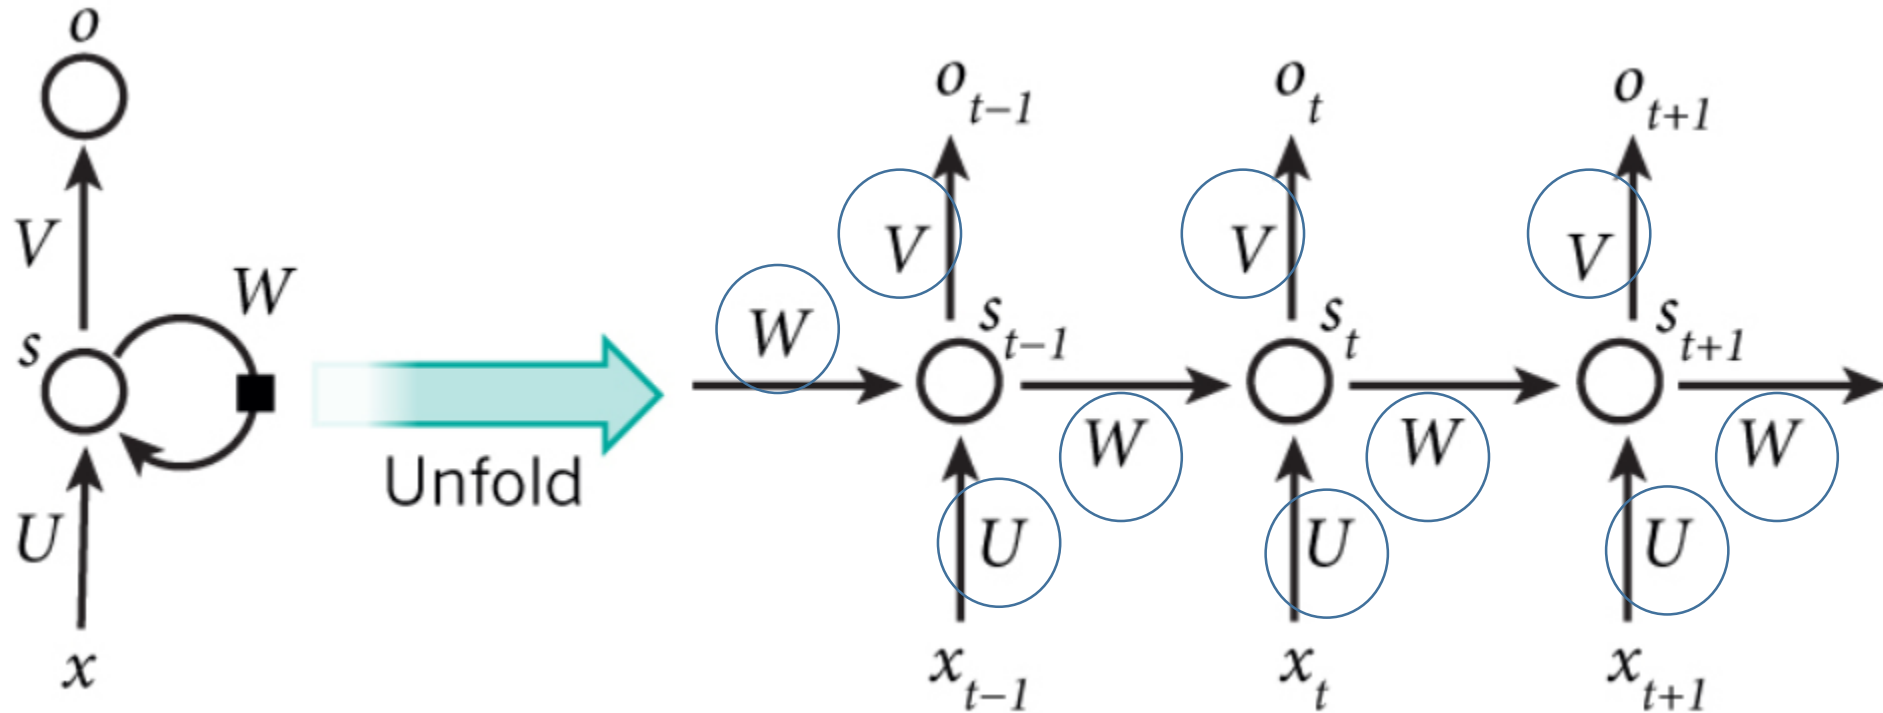


# Recurrent Neural Networks (RNN)

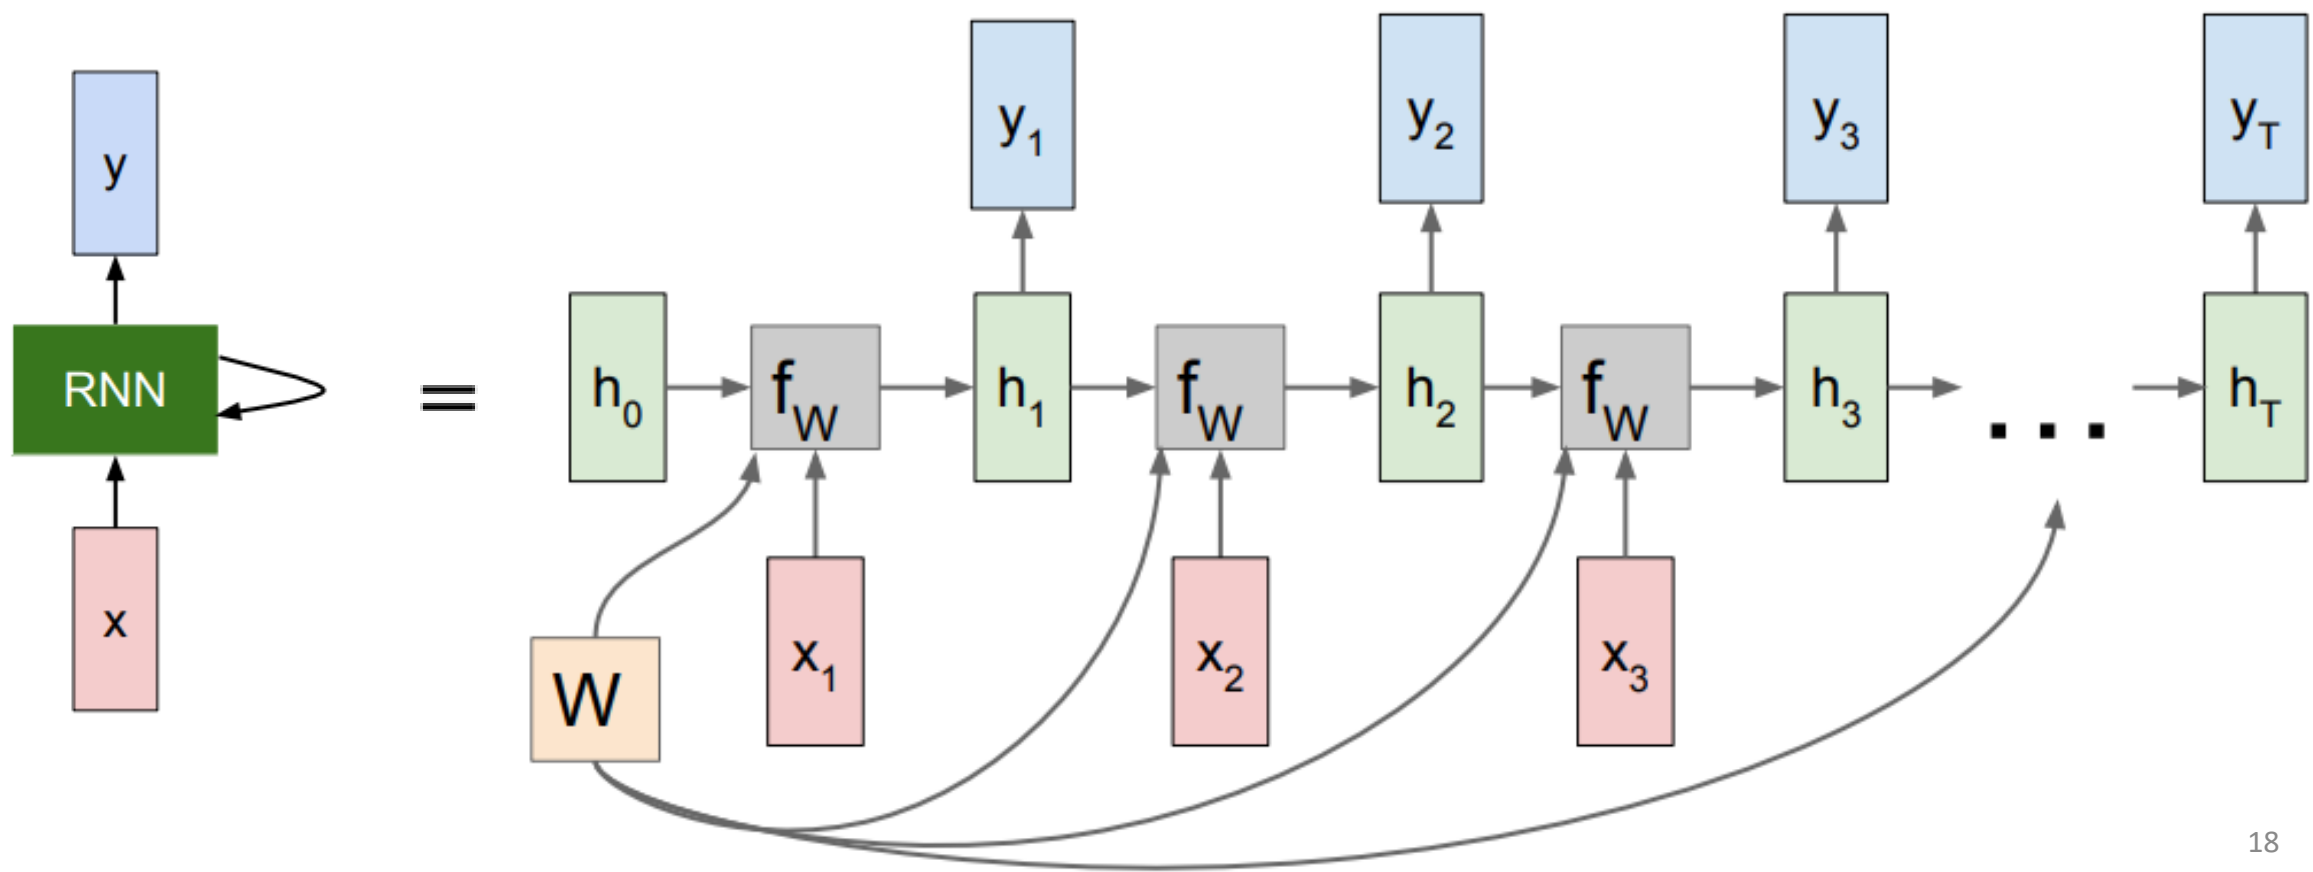




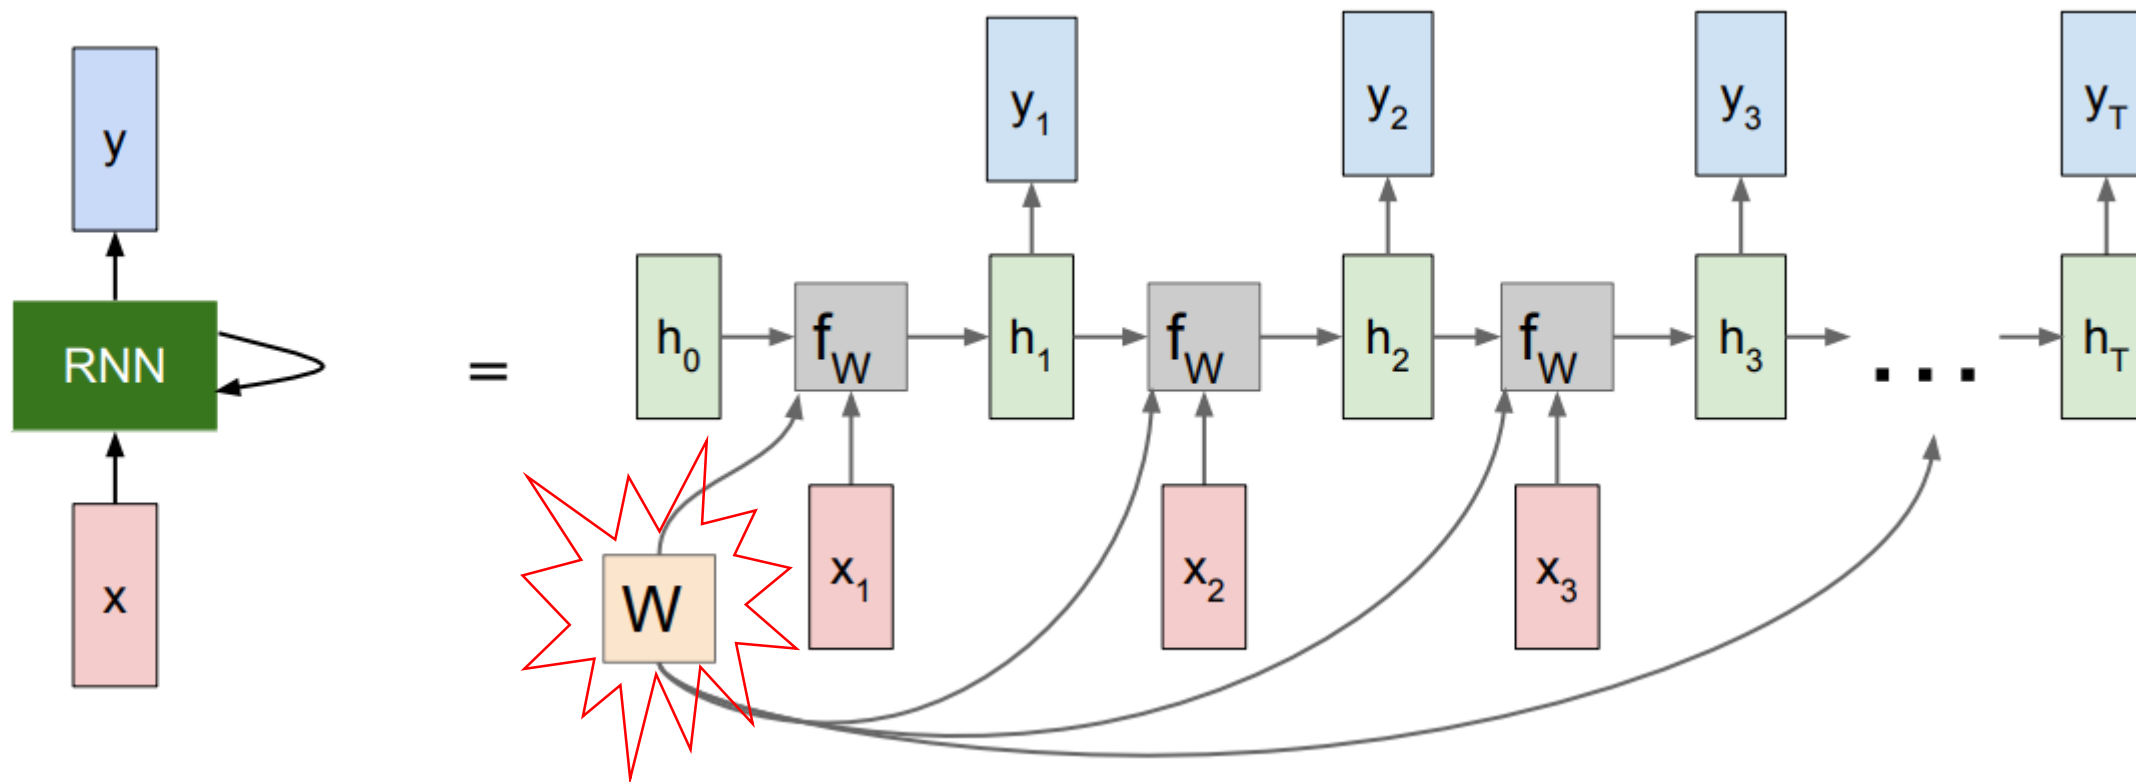
# Recurrent Neural Networks (RNN)



# RNN Computational Graph

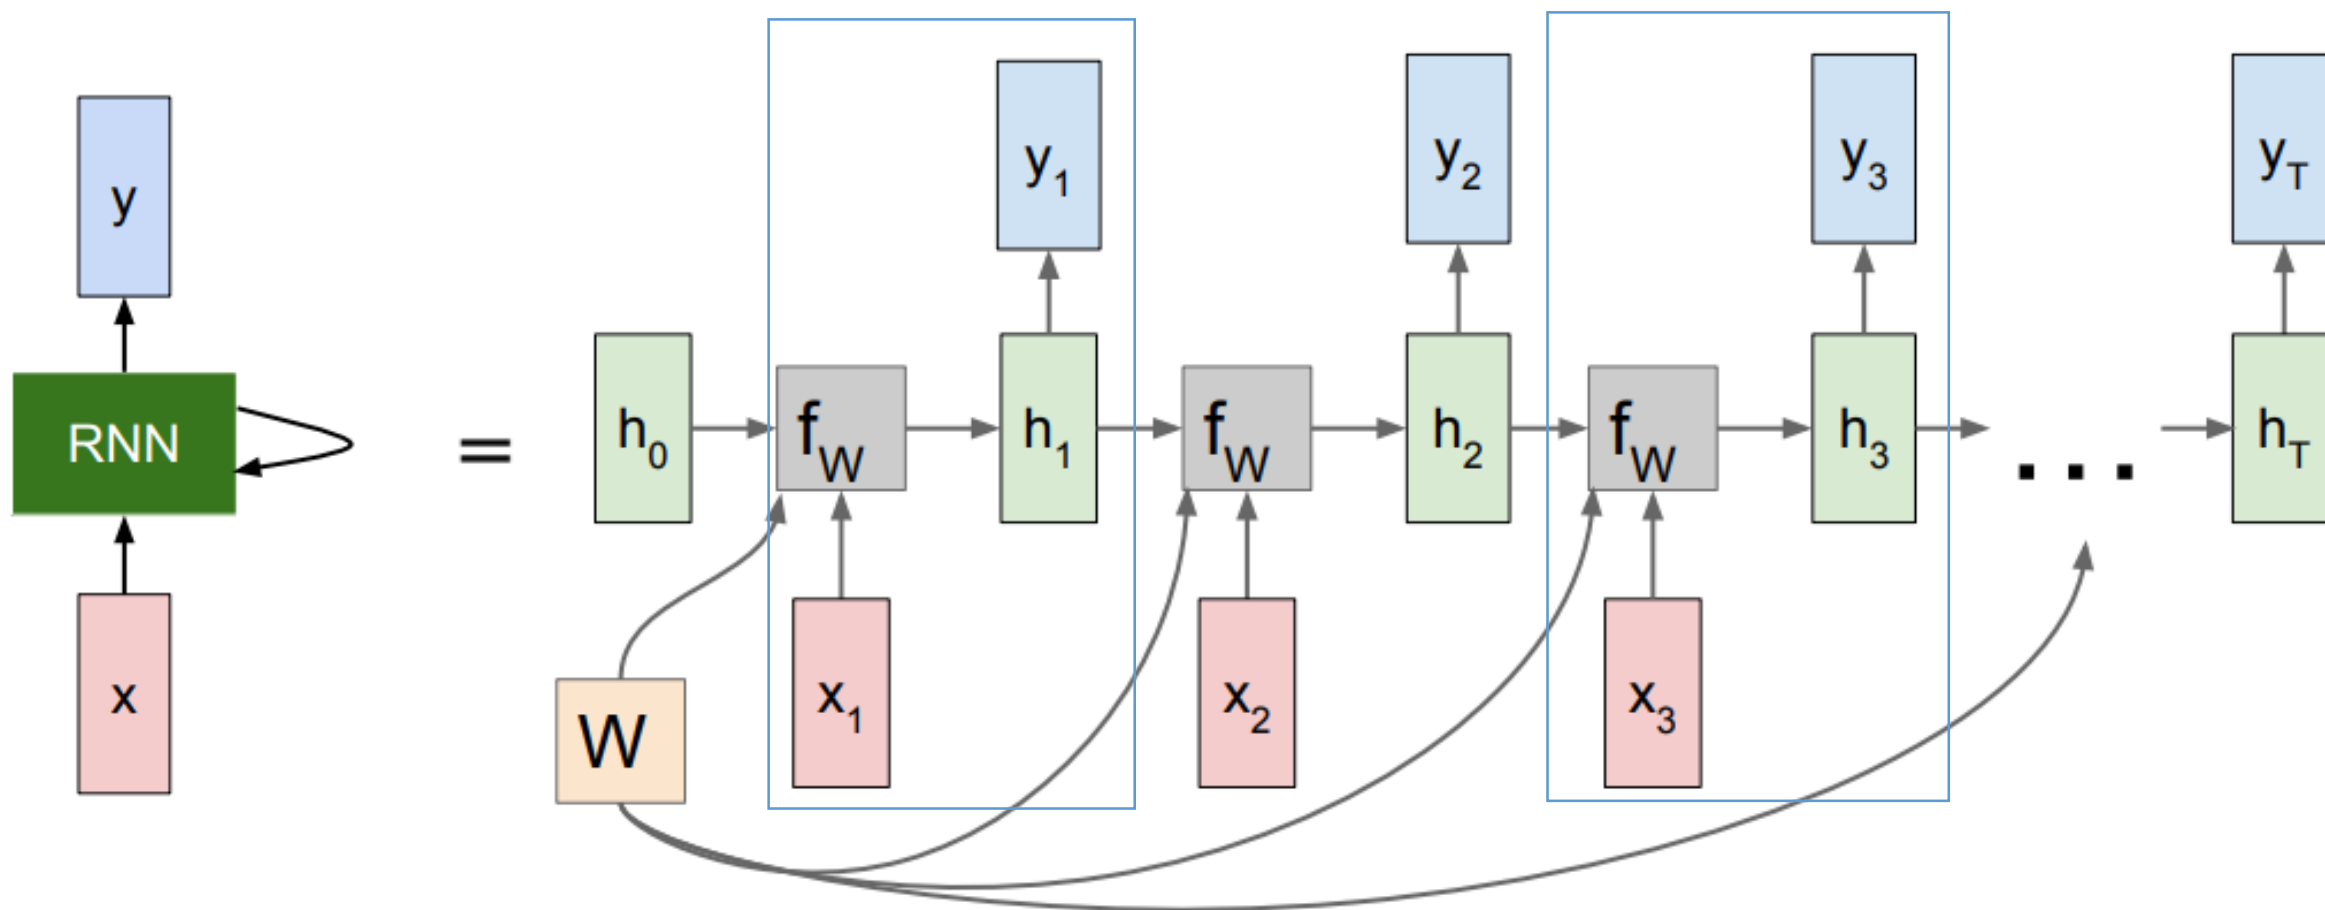


# RNN Computational Graph

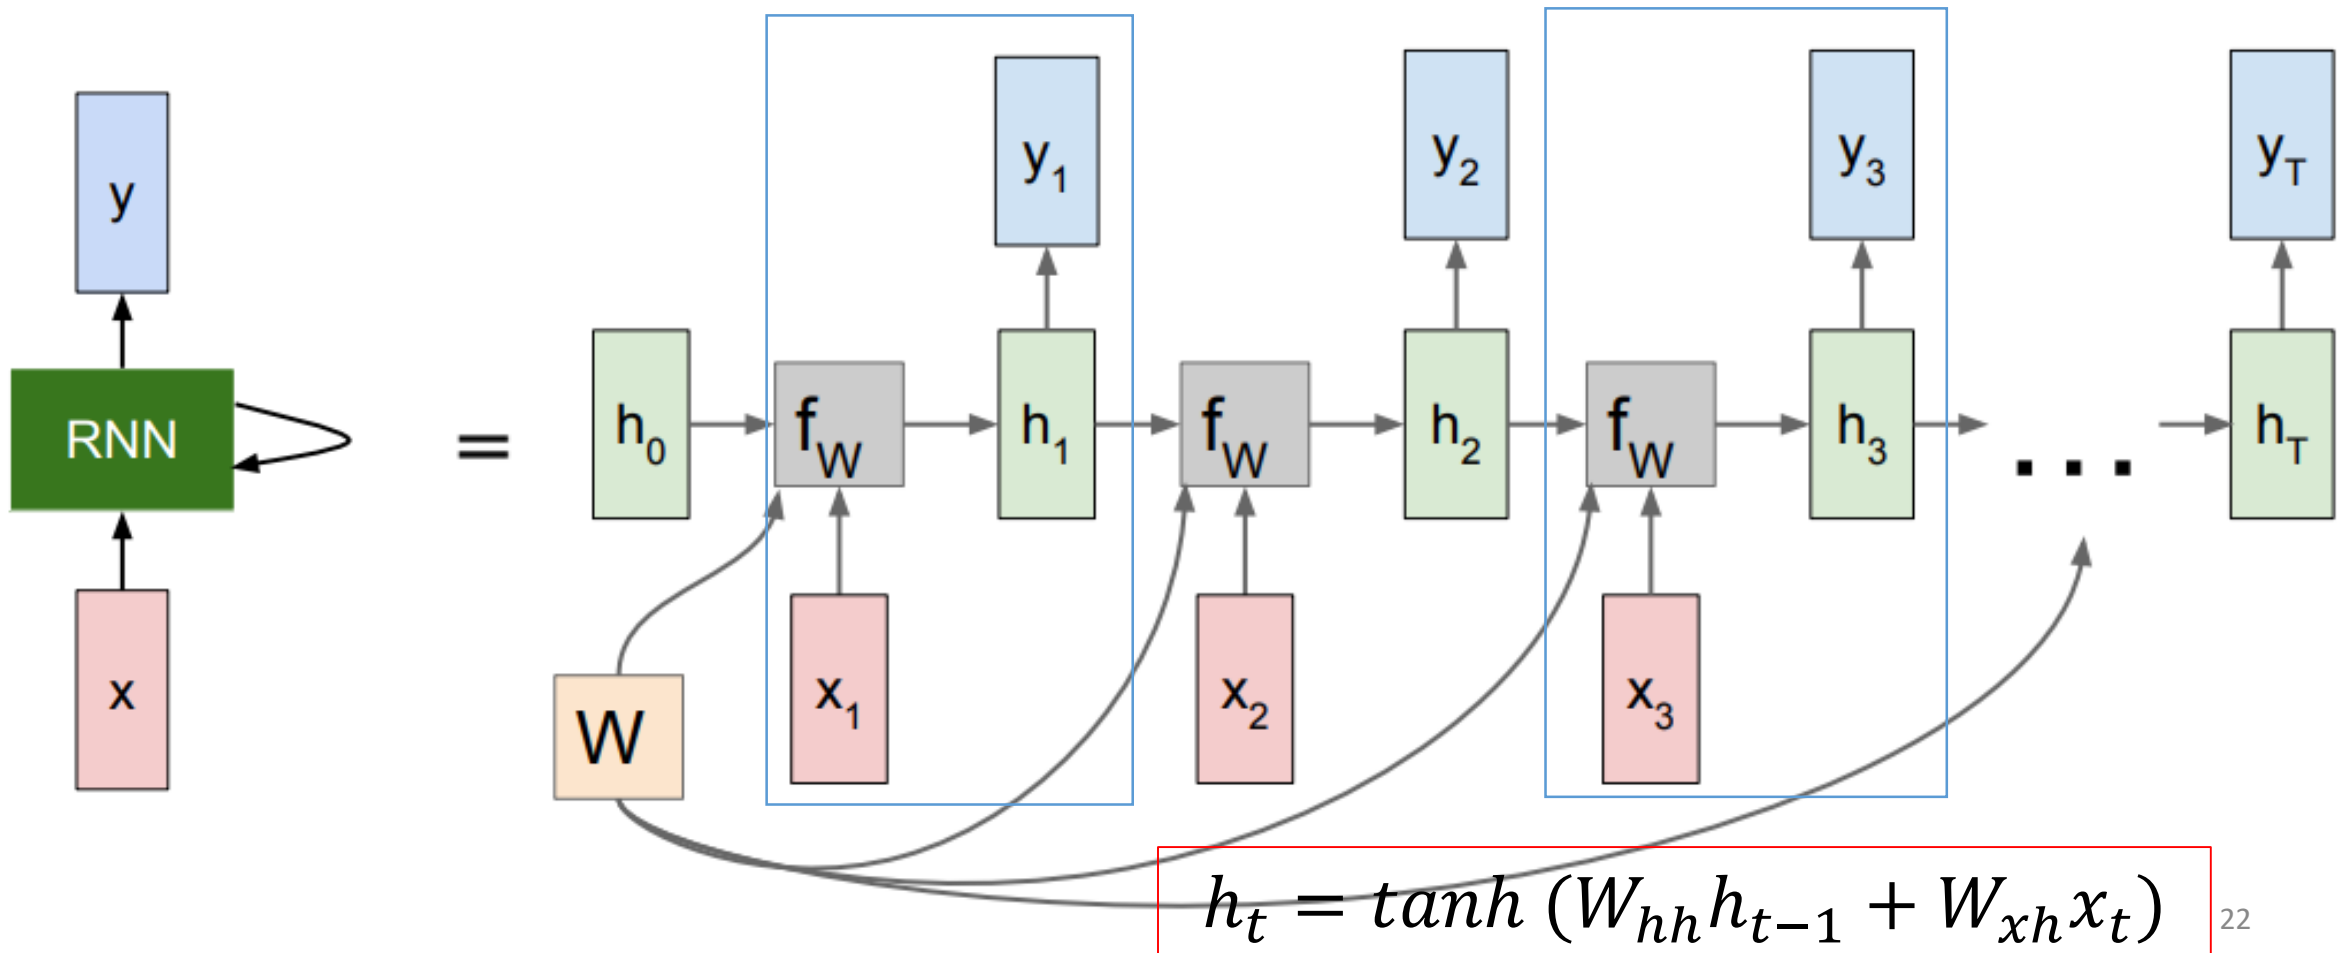


Do you notice any problem?

# RNN Computational Graph



# RNN Computational Graph



# Is it parallelizable?

- No! Note that for having  $h_t$  we first need  $h_{t-1}$ .

$$h_t = \tanh (W_{hh}h_{t-1} + W_{xh}x_t)$$

# Is it parallelizable?

- No! Note that for having  $h_t$  we first need  $h_{t-1}$ .
- How to solve it??

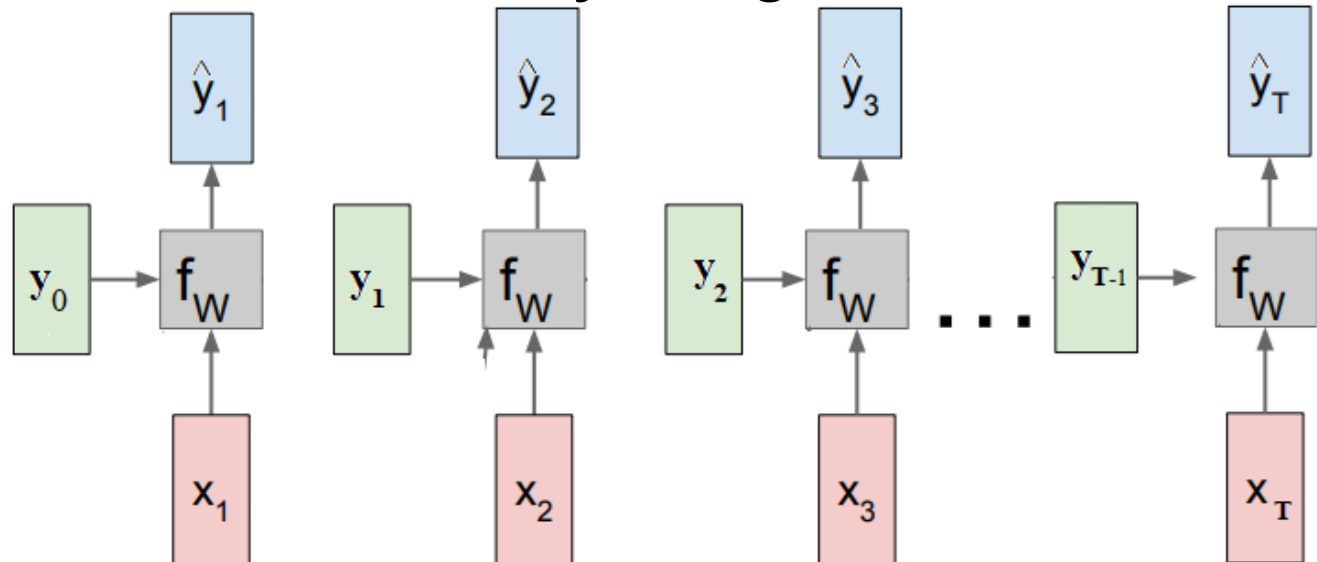


# Is it parallelizable?

- No! Note that for having  $h_t$  we first need  $h_{t-1}$ .
- How to solve it??
  - For some architectures we have *teacher forcing*


# Is it parallelizable?

- No! Note that for having  $h_t$  we first need  $h_{t-1}$ .
- How to solve it??
  - For some architectures we have *teacher forcing*




# Before continuing...

- Translate a phrase from English to French.
- Is this architecture suitable for this problem?



# Before continuing...

- Translate a phrase from English to French.
- Is this architecture suitable for this problem?



- Ans: No! Sentences might have different amount of words.  
We need to know the entire sentence before translating!

# Types

One to one




Image Classification

One to many





Image Captioning


Many to one



Sentiment Classification


# Types

Many to many




Machine Translation

Many to many




Video Classification on Frame Level

For each output...



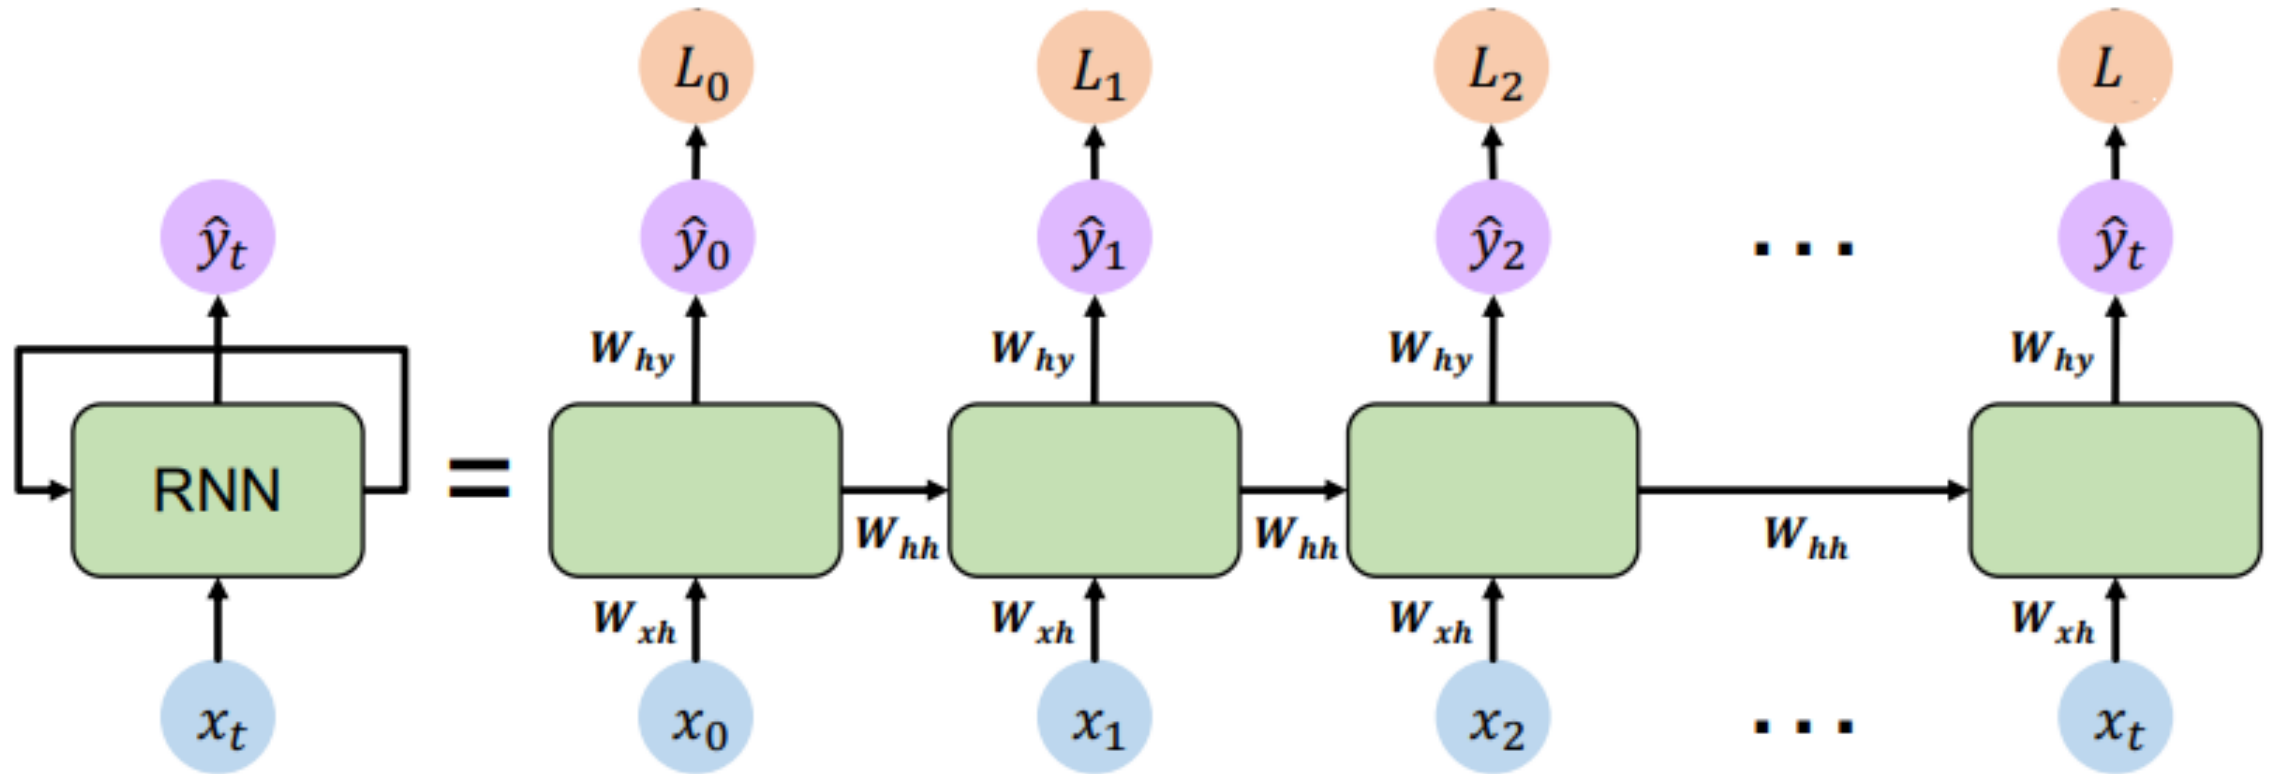
For each output...



$$E_0(y_0, \hat{y}_0) = -y_0 \log(\hat{y}_0)$$

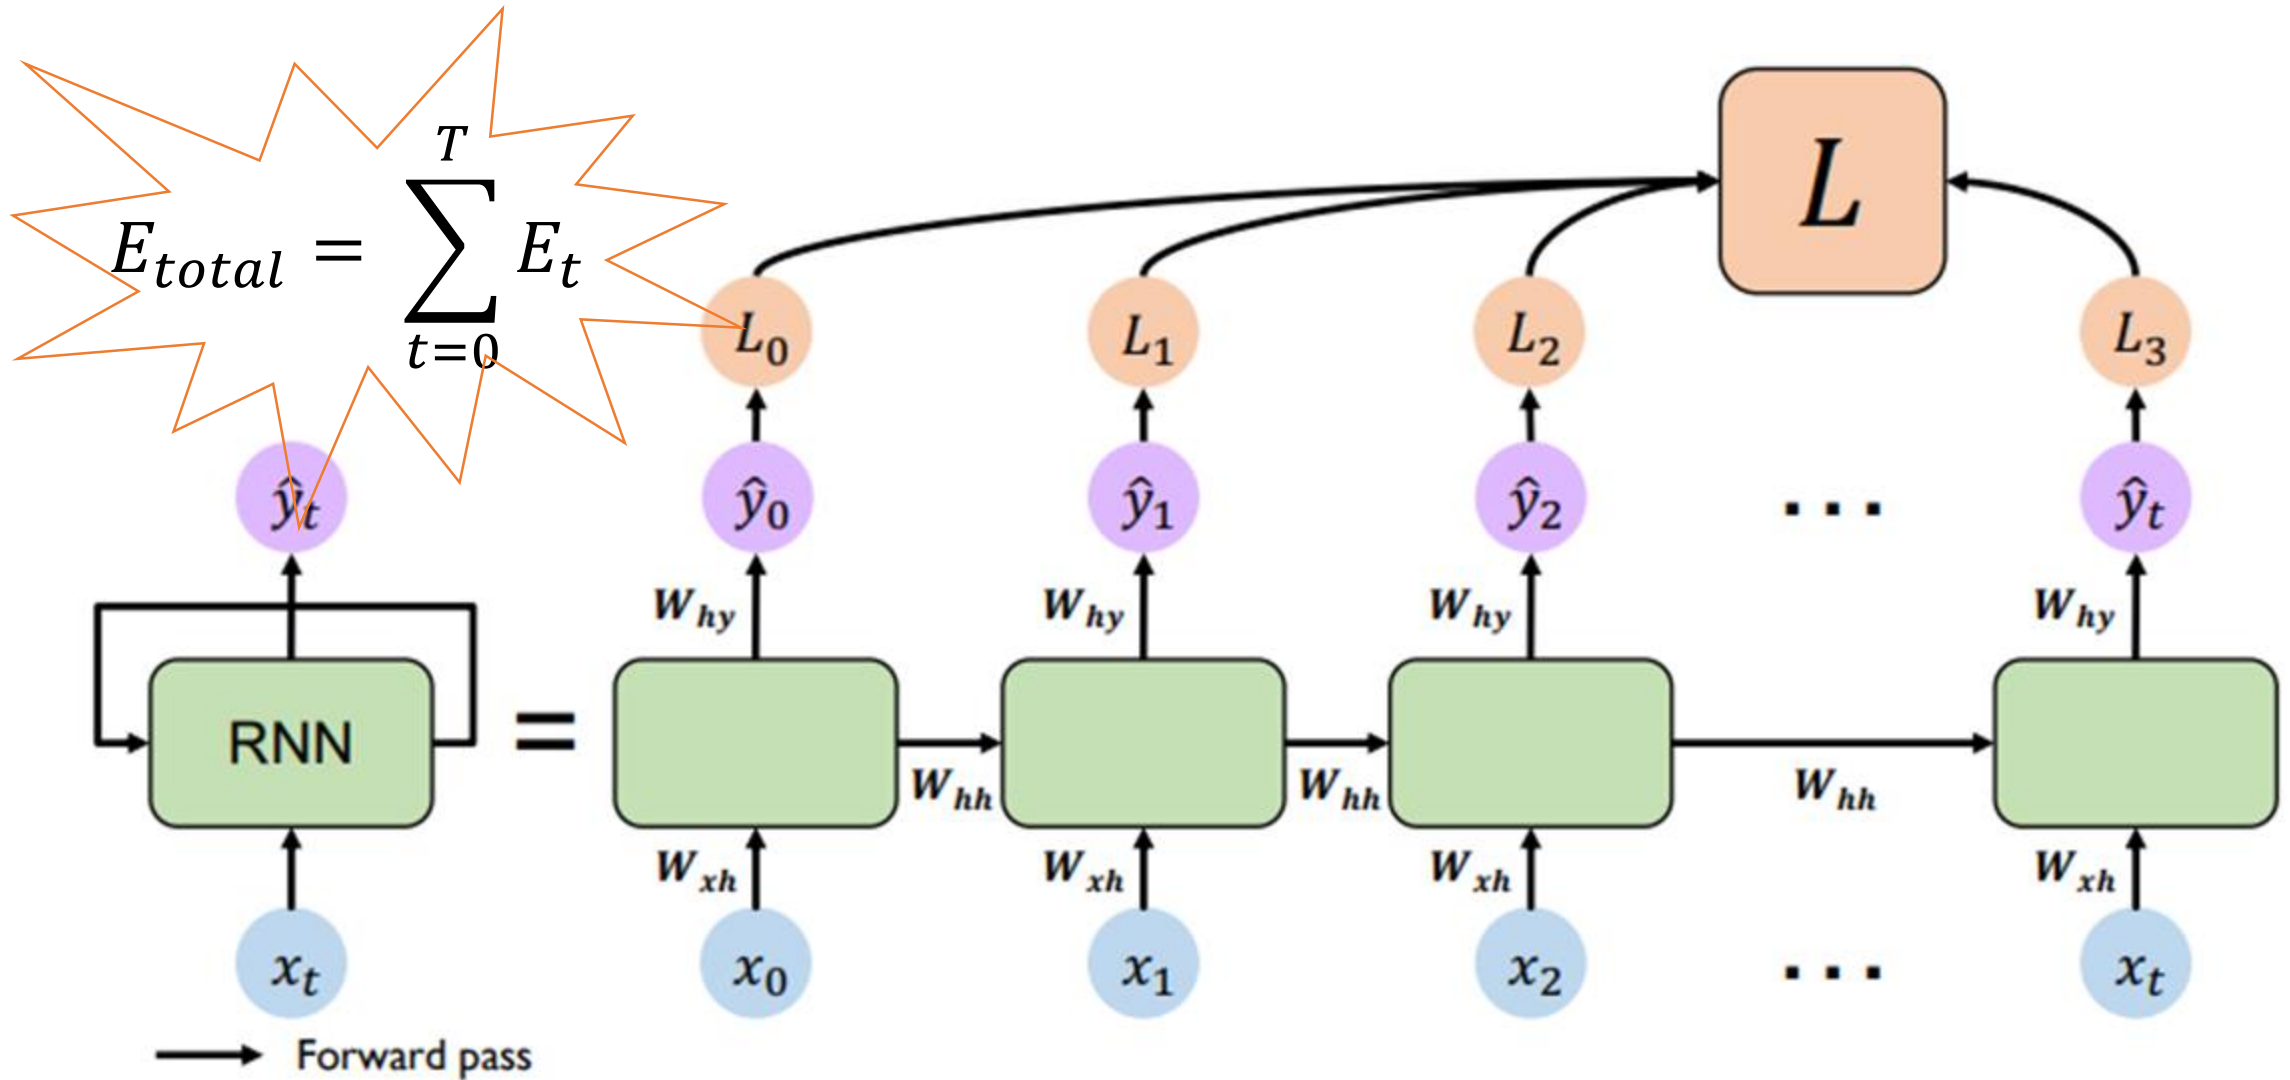


For each output...




$$E_t(y_t, \hat{y}_t) = -y_t \log(\hat{y}_t)$$


For each output...



# Backpropagation through time (BPTT)



# Backpropagation through time (BPTT)




- We must backpropagate through  $W$ ,  $V$  and  $U$
- Let's assume we are in the third cell...

$$\begin{aligned}
 \frac{\partial E_3}{\partial V} &= \frac{\partial E_3}{\partial \hat{y}_3} \frac{\partial \hat{y}_3}{\partial V} \\
 &= \frac{\partial E_3}{\partial \hat{y}_3} \frac{\partial \hat{y}_3}{\partial z_3} \frac{\partial z_3}{\partial V} \longrightarrow z_3 = V s_3 \longrightarrow s_3 = \tanh(U x_t + W s_2) \\
 &= (\hat{y}_3 - y_3) \otimes s_3
 \end{aligned}$$

For  $V$  it depends only on the values in cell number 3! But now let's see  $W$  and  $U$ .

# Backpropagation through time (BPTT)

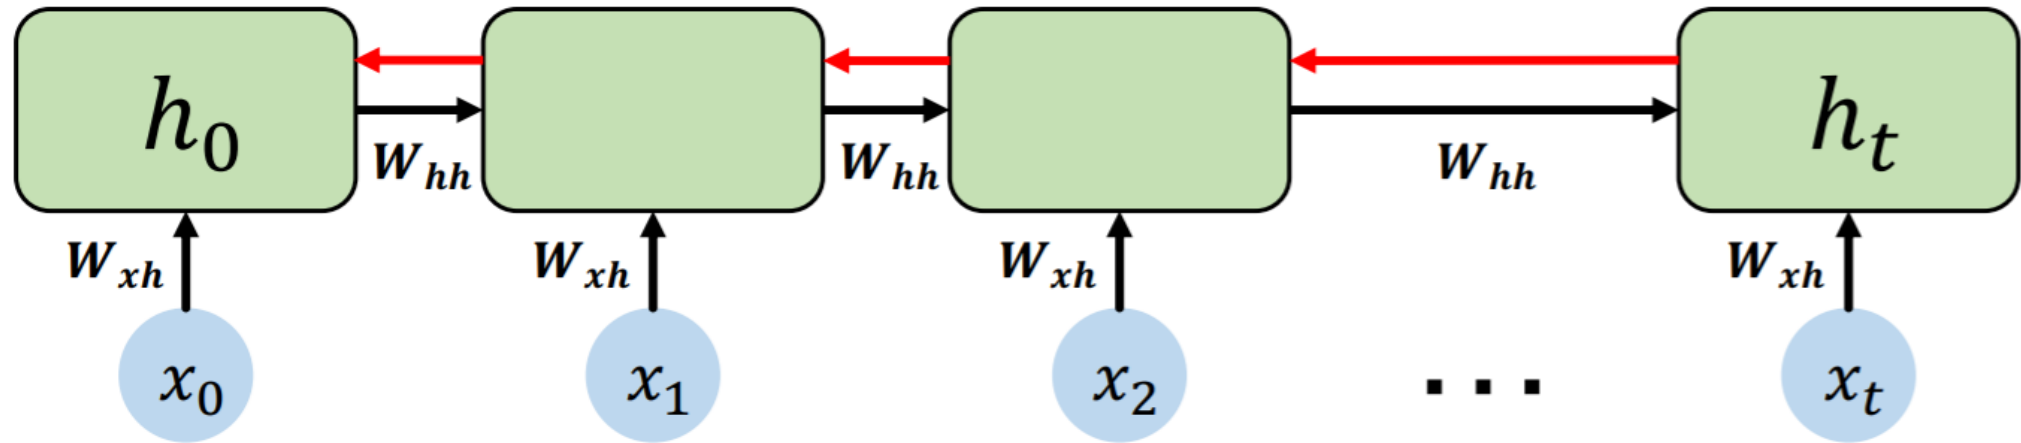


$$\frac{\partial E_3}{\partial W} = \frac{\partial E_3}{\partial \hat{y}_3} \frac{\partial \hat{y}_3}{\partial s_3} \frac{\partial s_3}{\partial W} \longrightarrow s_3 = \tanh(Ux_t + Ws_2)$$

$$\frac{\partial E_3}{\partial W} = \sum_{k=0}^3 \frac{\partial E_3}{\partial \hat{y}_3} \frac{\partial \hat{y}_3}{\partial s_3} \frac{\partial s_3}{\partial s_k} \frac{\partial s_k}{\partial W}$$

The gradient will depend in all the previous cells!  
It's the same situation for U!

# Vanishing/Exploding Gradients Problem




Computing the BPTT algorithm involves the multiplication of many factors of the derivative of  $W$  ( $W_{hh}$ ) and the activation function ( $\tanh$ ). Therefore it's possible to have some gradient problems.

# Vanishing/Exploding Gradients Problem

Many of the values  $> 1$   
**Exploding gradients**



- 1. Gradient clipping:**  
Scale the gradient if it's bigger than a threshold.
- 2. Truncated BPTT**




# Truncated backpropagation through time (TBTT)


In long RNN's BPTT can be time consuming because it has to go through a lot of cells. So some researches are just backpropagating through a pre-established number of cells.



# TBTT



# TBTT



# Vanishing/Exploding Gradients Problem

Many of the values  $< 1$


**Vanishing gradients**




**Bias network to short-term dependencies.**

1. Change Activation function.
2. Weight Initialization
3. Change RNN architecture

“The clouds are in the \_\_\_\_”




“I grew up in France, ... and I I speak fluent\_\_\_\_”



# Vanishing/Exploding Gradients Problem

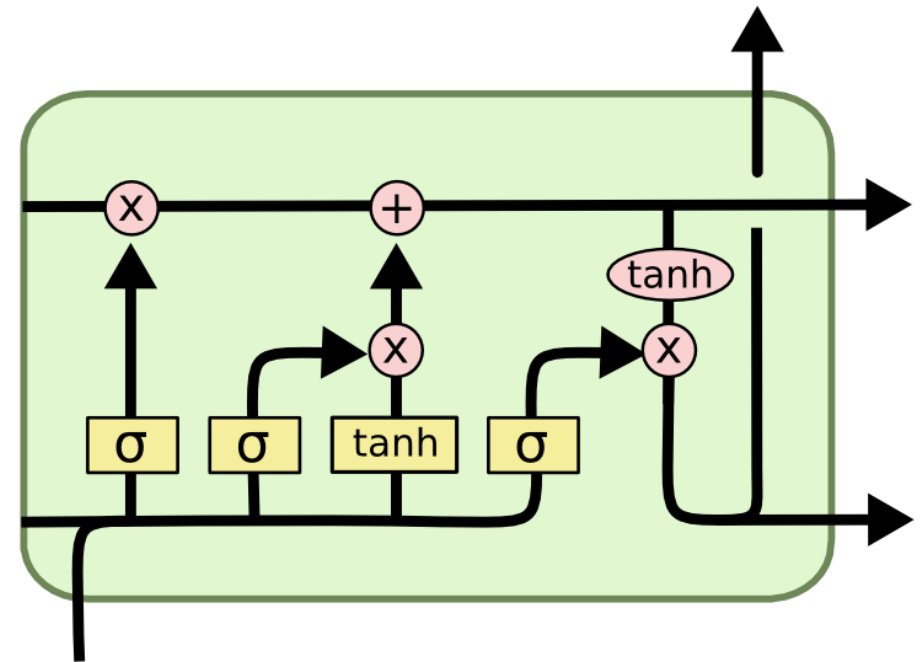

## Activation Function



## Weight Initialization

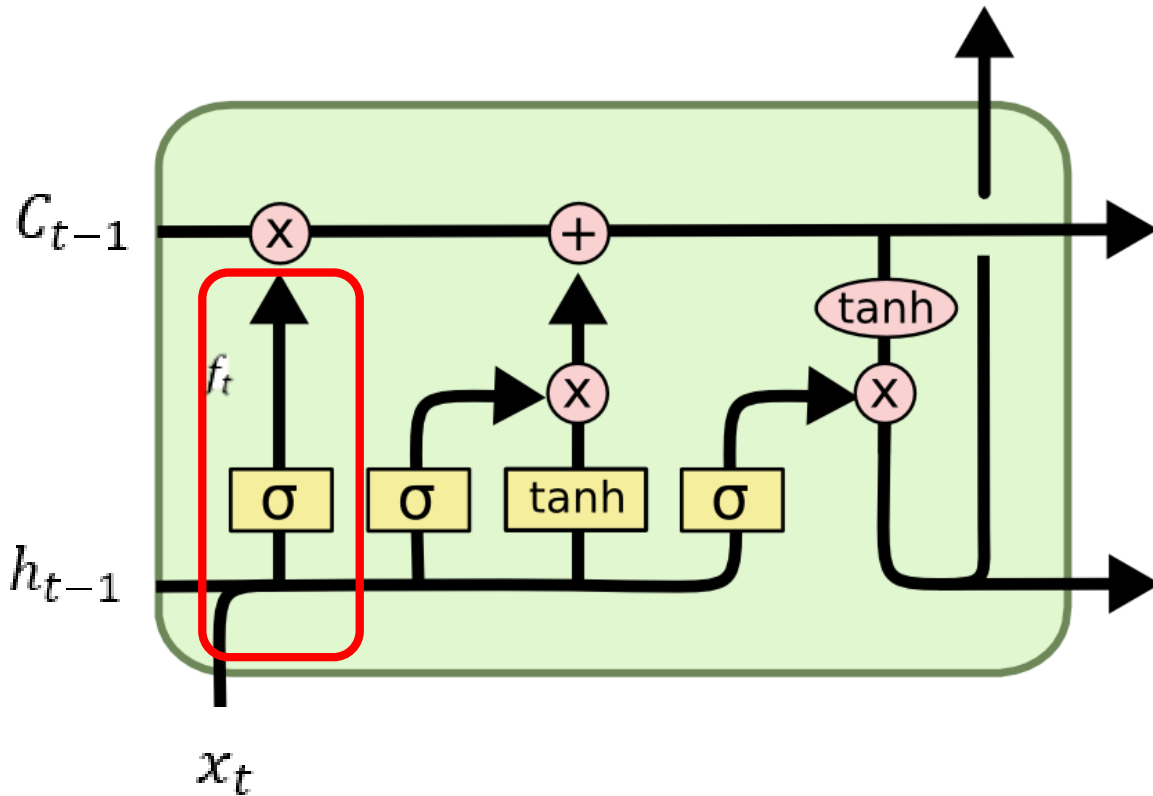
$$I_n = \begin{pmatrix} 1 & 0 & 0 & \dots & 0 \\ 0 & 1 & 0 & \dots & 0 \\ 0 & 0 & 1 & \dots & 0 \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & 0 & \dots & 1 \end{pmatrix}$$

# Long-Short Term Memory (LSTM)



# Long-Short Term Memory (LSTM)

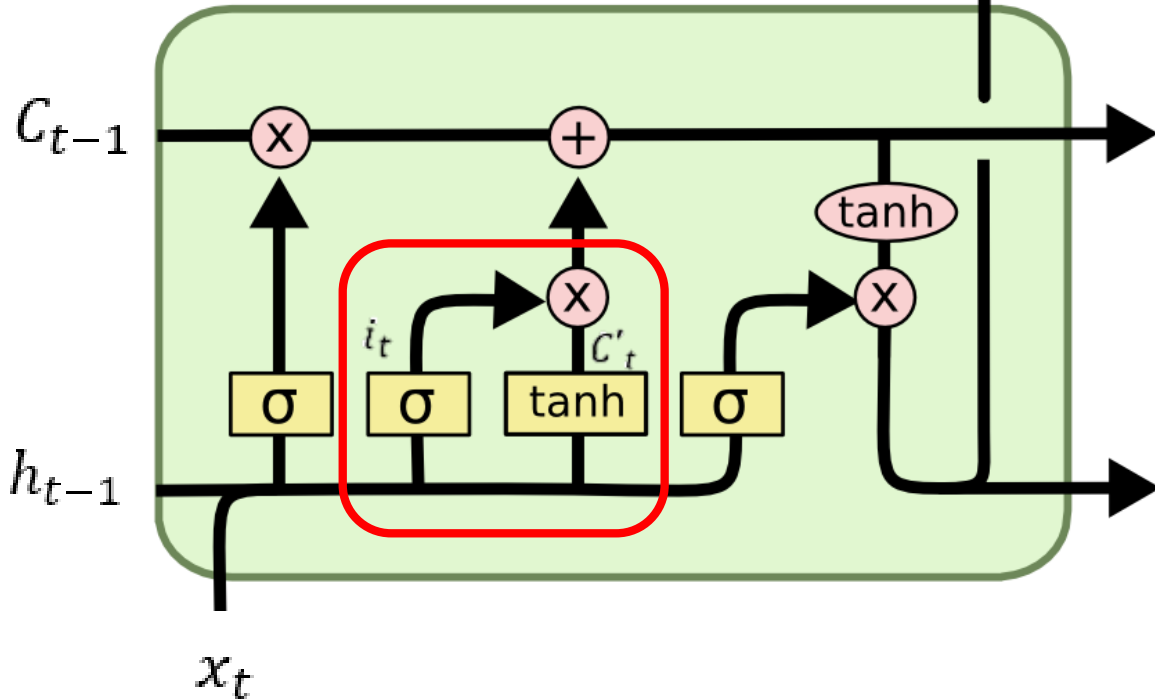
**Forget Gate – What to forget about the input and the hidden state of the previous cell.**



$$f_t = \sigma(W_i[h_{t-1}, x_t] + b_f)$$

# Long-Short Term Memory (LSTM)

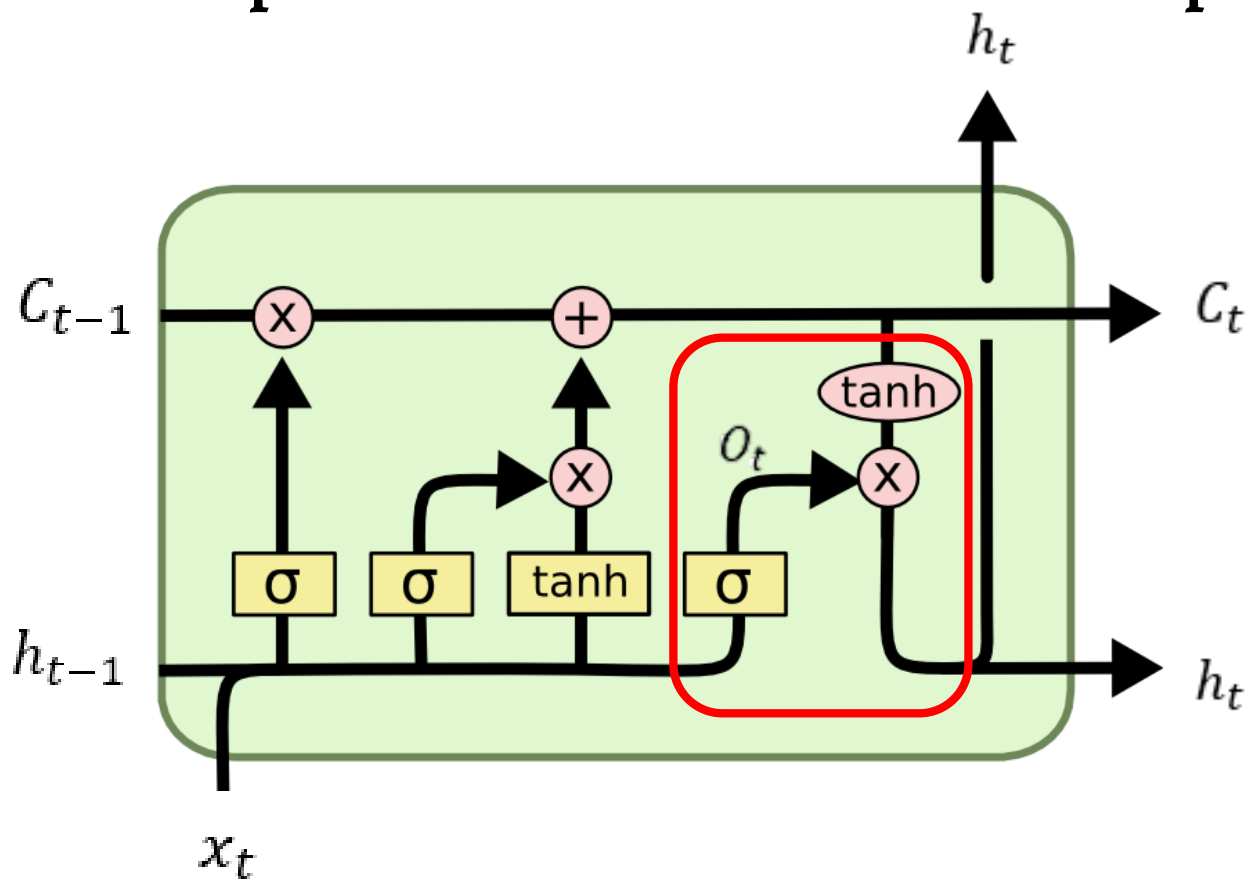
**Input Gate – What to write about the input and the hidden state of the previous cell into the cell state**



$$i_t = \sigma(W_i[h_{t-1}, x_t] + b_i)$$
$$\tilde{C}_t = \tanh(W_c[h_{t-1}, x_t] + b_c)$$

# Long-Short Term Memory (LSTM)


**Output Gate – What to show about the input and the cell state to the next cell.**



$$o_t = \sigma(W_o [h_{t-1}, x_t] + b_o)$$

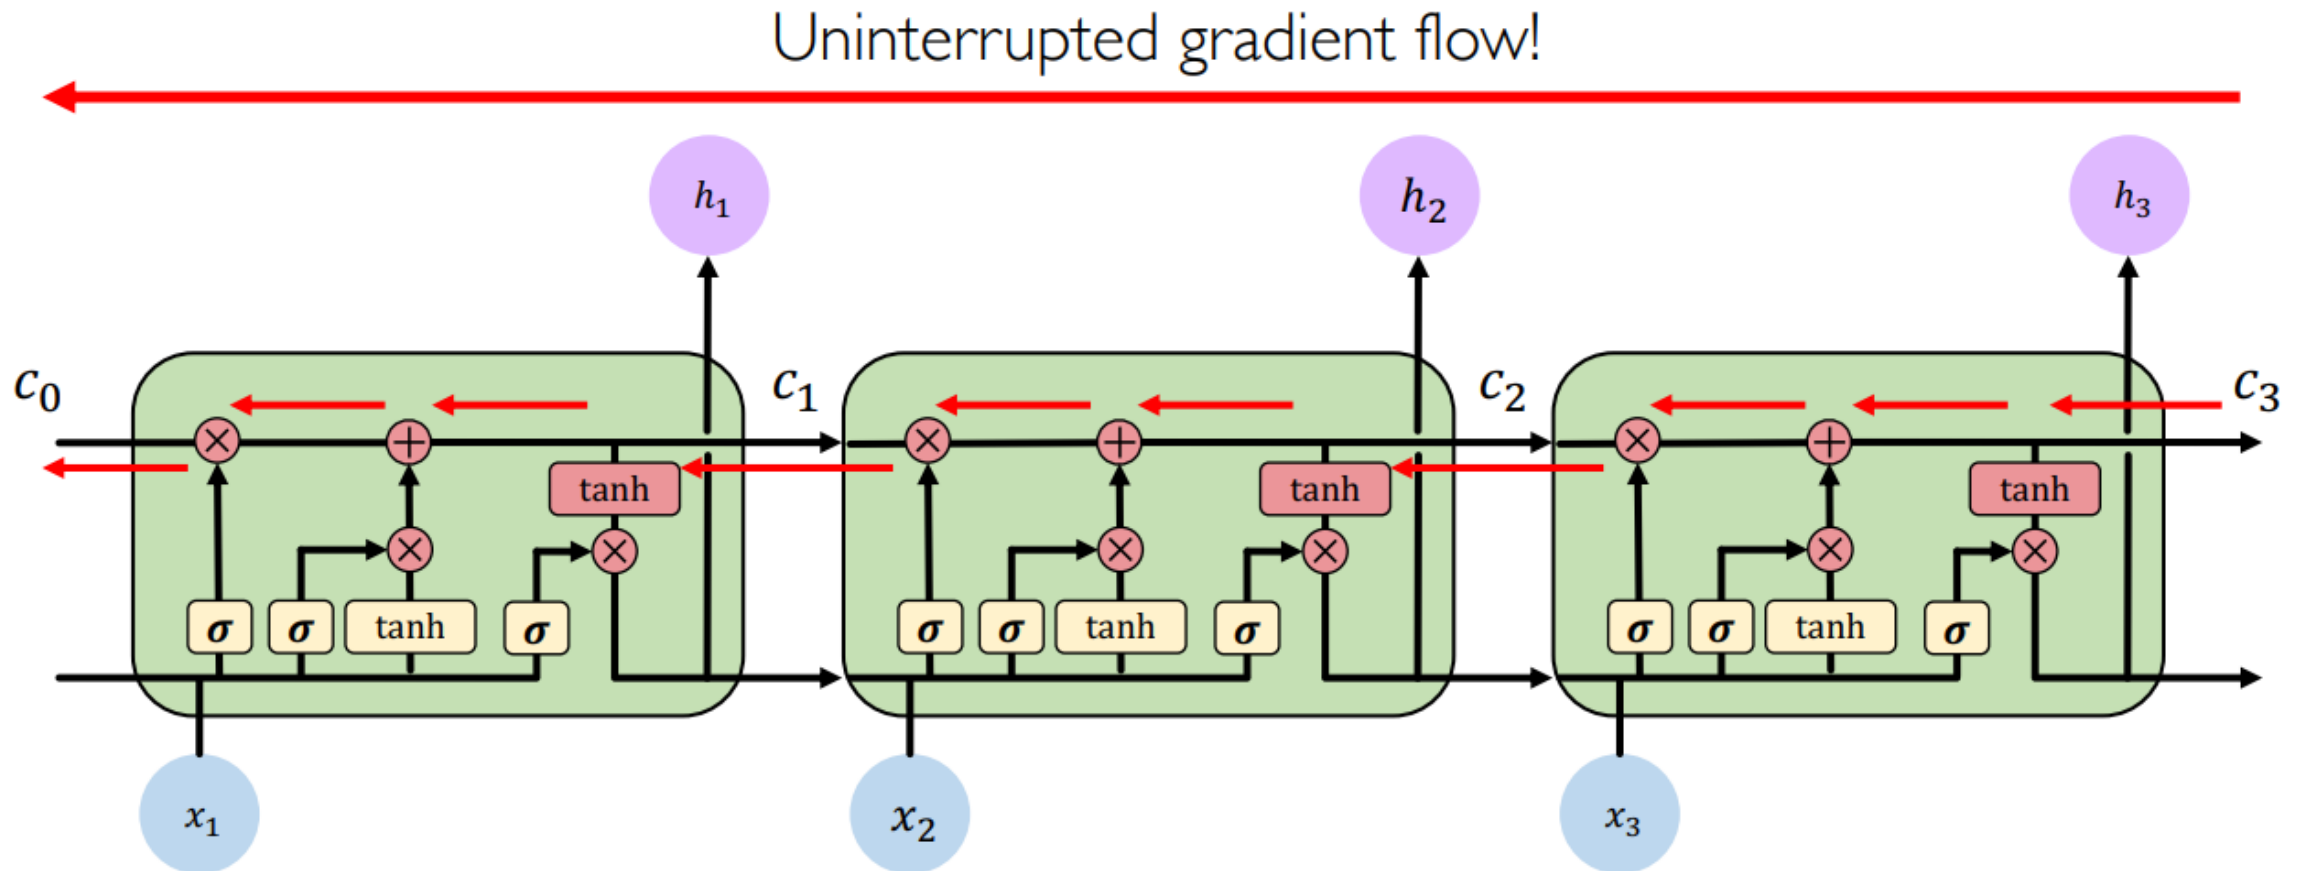


# Long-Short Term Memory (LSTM)




$$c_t = f \odot c_{t-1} + i \odot c'_t$$
$$h_t = o \odot \tanh(c_t)$$

# Long-Short Term Memory (LSTM)




# Long-Short Term Memory (LSTM)

Backward process only dependent with the forget gate!!!




# Long-Short Term Memory (LSTM)

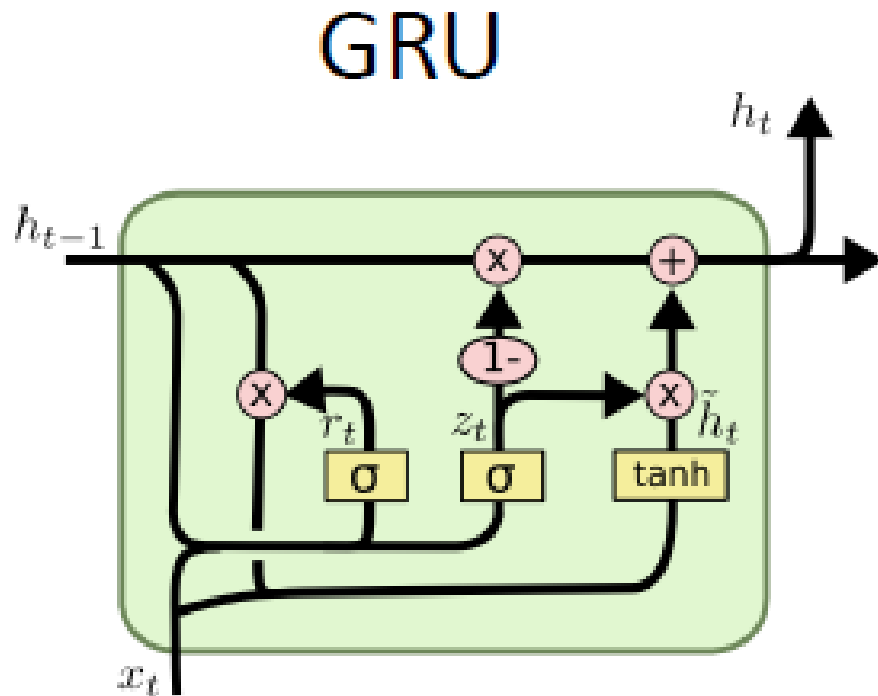
Uninterrupted gradient flow!



Similar to RESNET!



# Gradient Recurrent Unit(GRU)



$$r_t = \sigma(W_{xr}x_t + W_{hr}h_{t-1} + b_r)$$


$$z_t = \sigma(W_{xz}x_t + W_{hz}h_{t-1} + b_z)$$

$$\tilde{h}_t = \tanh(W_{xh}x_t + W_{hh}(r_t \odot h_{t-1}) + b_h)$$

$$h_t = z_t \odot h_{t-1} + (1 - z_t) \odot \tilde{h}_t$$

# Gradient Recurrent Unit (GRU)


**Update Gate – Determine how much of past information needs to pass to next step**



$$z_t = \sigma(W_z x_t + U_z h_{t-1})$$

# Gradient Recurrent Unit (GRU)


**Reset Gate – Decide how much of past information forget**



$$r_t = \sigma(W_r x_t + U_r h_{t-1})$$

# Gradient Recurrent Unit (GRU)

**Current Memory Content**– Use the reset gate to store relevant information from the past




$$\tilde{h}_t = \tanh(Wx_t + r_t \odot Uh_{t-1})$$




# Gradient Recurrent Unit (GRU)

**Final Current memory at time step – Holds information for the current unit.**



$$h_t = z_t \odot h_{t-1} + (1 - z_t) \odot \tilde{h}_t$$


# GRU: Current Memory Content



The new memory content will use the reset gate to store the **relevant** information from the past.

$$\tilde{h}_t = \tanh(Wx_t + r_t \odot Uh_{t-1})$$

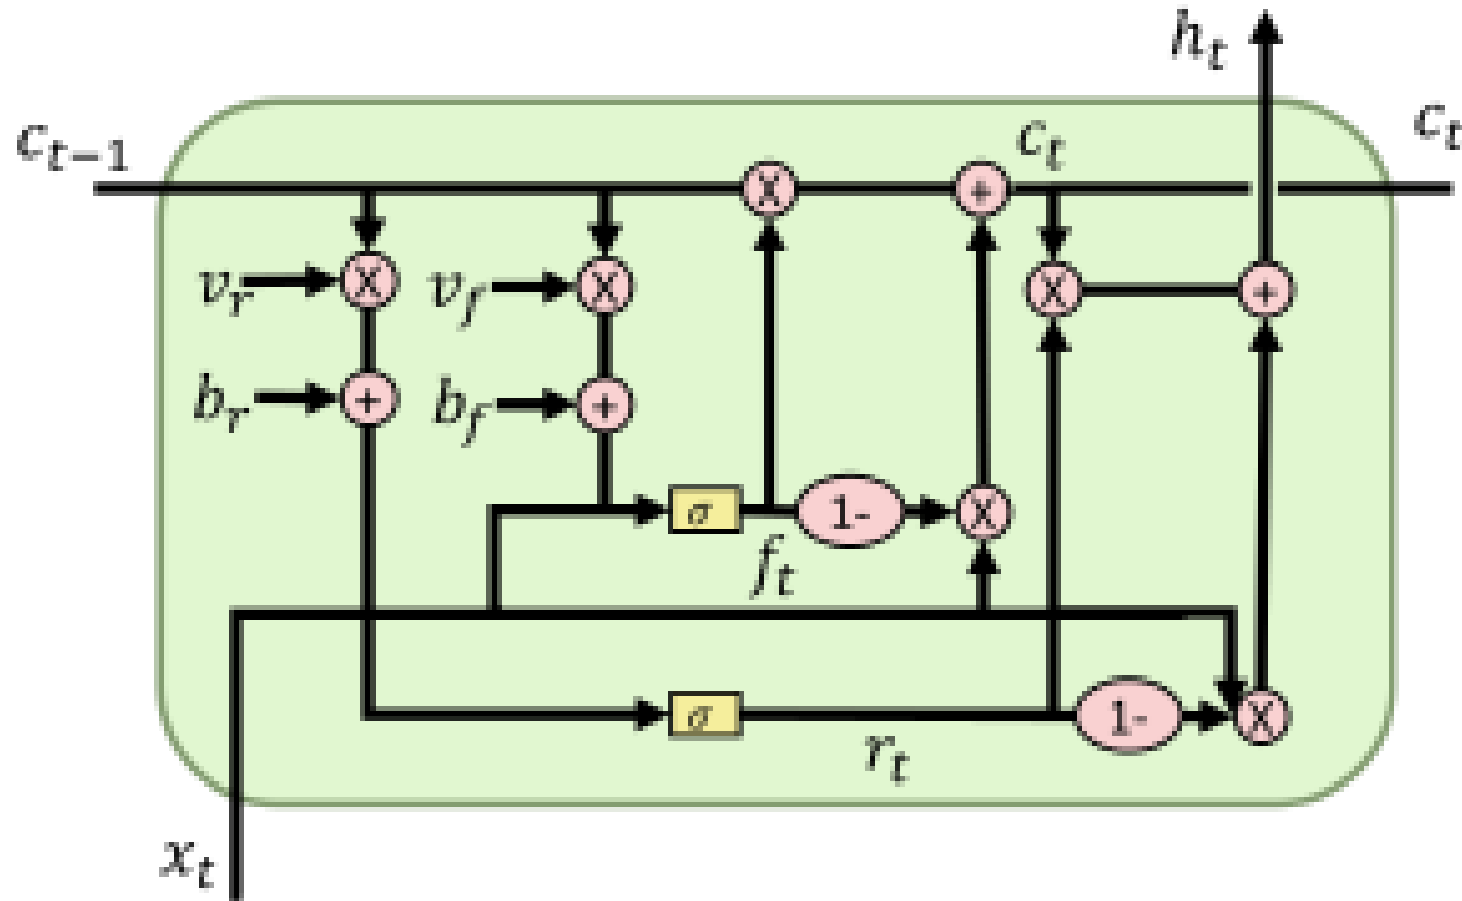
# GRU: Final Current memory at time step



Holds information for the  
current unit and passes it  
down to the network.

$$h_t = z_t \odot h_{t-1} + (1 - z_t) \odot \tilde{h}_t$$

# Simple Recurrent Unit SRU)




$$\mathbf{f}_t = \sigma(\mathbf{W}_f \mathbf{x}_t + \mathbf{v}_f \odot \mathbf{c}_{t-1} + \mathbf{b}_f)$$

$$\mathbf{c}_t = \mathbf{f}_t \odot \mathbf{c}_{t-1} + (1 - \mathbf{f}_t) \odot (\mathbf{W} \mathbf{x}_t)$$

$$\mathbf{r}_t = \sigma(\mathbf{W}_r \mathbf{x}_t + \mathbf{v}_r \odot \mathbf{c}_{t-1} + \mathbf{b}_r)$$

$$\mathbf{h}_t = \mathbf{r}_t \odot \mathbf{c}_t + (1 - \mathbf{r}_t) \odot \mathbf{x}_t$$

# Simple Recurrent Unit(SRU)



**Light recurrence:**

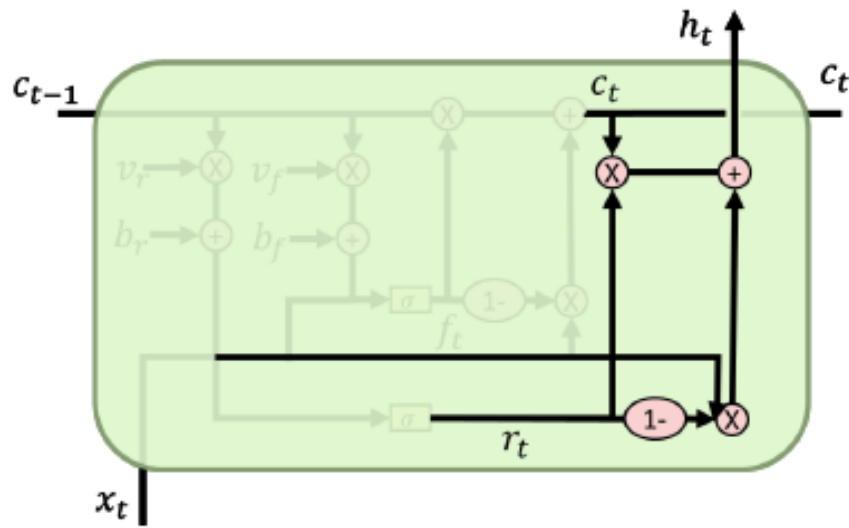
1. **Forget gate: Controls information flow**
2. **State vector: Adaptively average the previous state and the current c**

$$f_t = \sigma(W_f x_t + v_f \odot c_{t-1} + b_f)$$

$$c_t = f_t \odot c_{t-1} + (1 - f_t) \odot (W x_t)$$

Weighted average according to the forgot gate

# Simple Recurrent Unit(SRU)




$$r_t = \sigma(W_r x_t + v_r \odot c_{t-1} + b_r)$$

$$h_t = r_t \odot c_t + (1 - r_t) \odot x_t$$

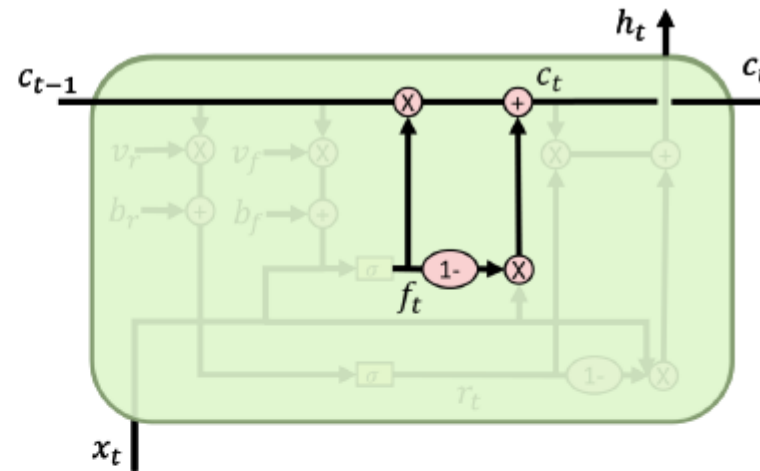
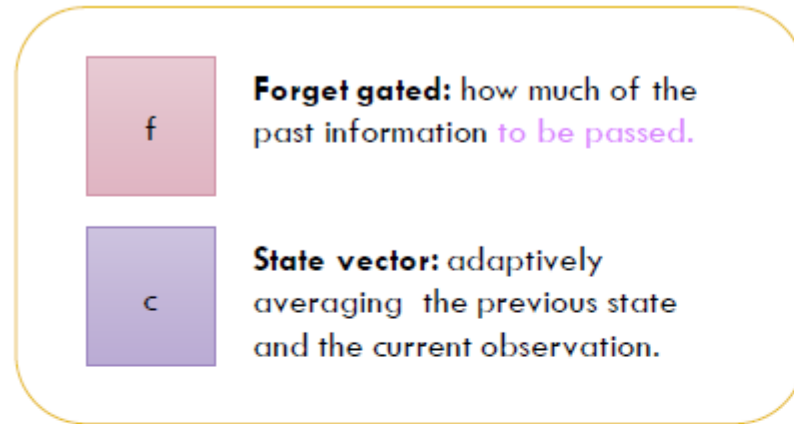
**Highway network:**

1. **Reset gate:** how much current information to be passed.
2. **Output vector:** Adaptively combine input and the state vector

# Variations – Bidirectional



# Simple recurrent unite (SRU)



$$f_t = \sigma(W_f x_t + v_f \odot c_{t-1} + b_f)$$

$$c_t = f_t \odot c_{t-1} + (1 - f_t) \odot (W x_t)$$

Weighted average according to the forgot gate



# Simple recurrent unite (SRU)

f

**Forget gated:** how much of the past information to be passed.

c

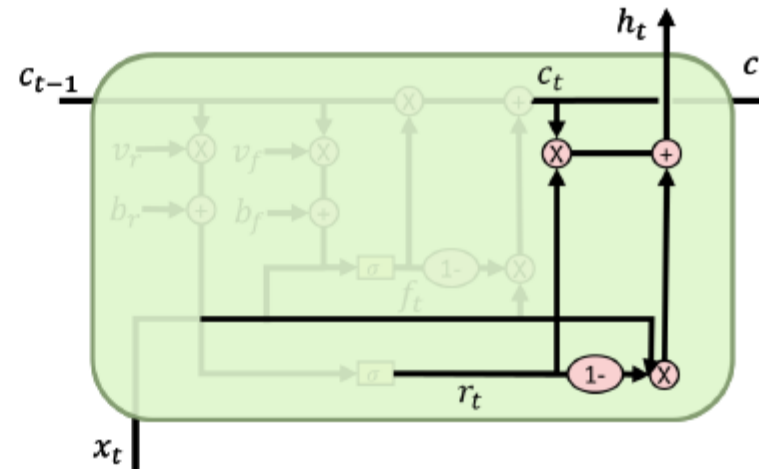
**State vector:** adaptively averaging the previous state and the current observation.

r

**Reset gated:** how much of the current information to be passed.

h



**Output vector:** adaptively combine the input and the state vector produced from the light recurrence.



$$r_t = \sigma(W_r x_t + v_r \odot c_{t-1} + b_r)$$



$$h_t = r_t \odot c_t + (1 - r_t) \odot x_t$$

# Variations – Multilayer



$$h_t^l = \tanh W^l \begin{pmatrix} h_t^{l-1} \\ h_{t-1}^l \end{pmatrix}$$


# Variations – LSTM Multilayer





$$\begin{pmatrix} i \\ f \\ o \\ g \end{pmatrix} = \begin{pmatrix} \text{sigm} \\ \text{sigm} \\ \text{sigm} \\ \text{tanh} \end{pmatrix} W^l \begin{pmatrix} h_t^{l-1} \\ h_{t-1}^l \end{pmatrix}$$
$$c_t^l = f \odot c_{t-1}^l + i \odot g$$
$$h_t^l = o \odot \tanh(c_t^l)$$

# Variations-

## Encoder-Decoder

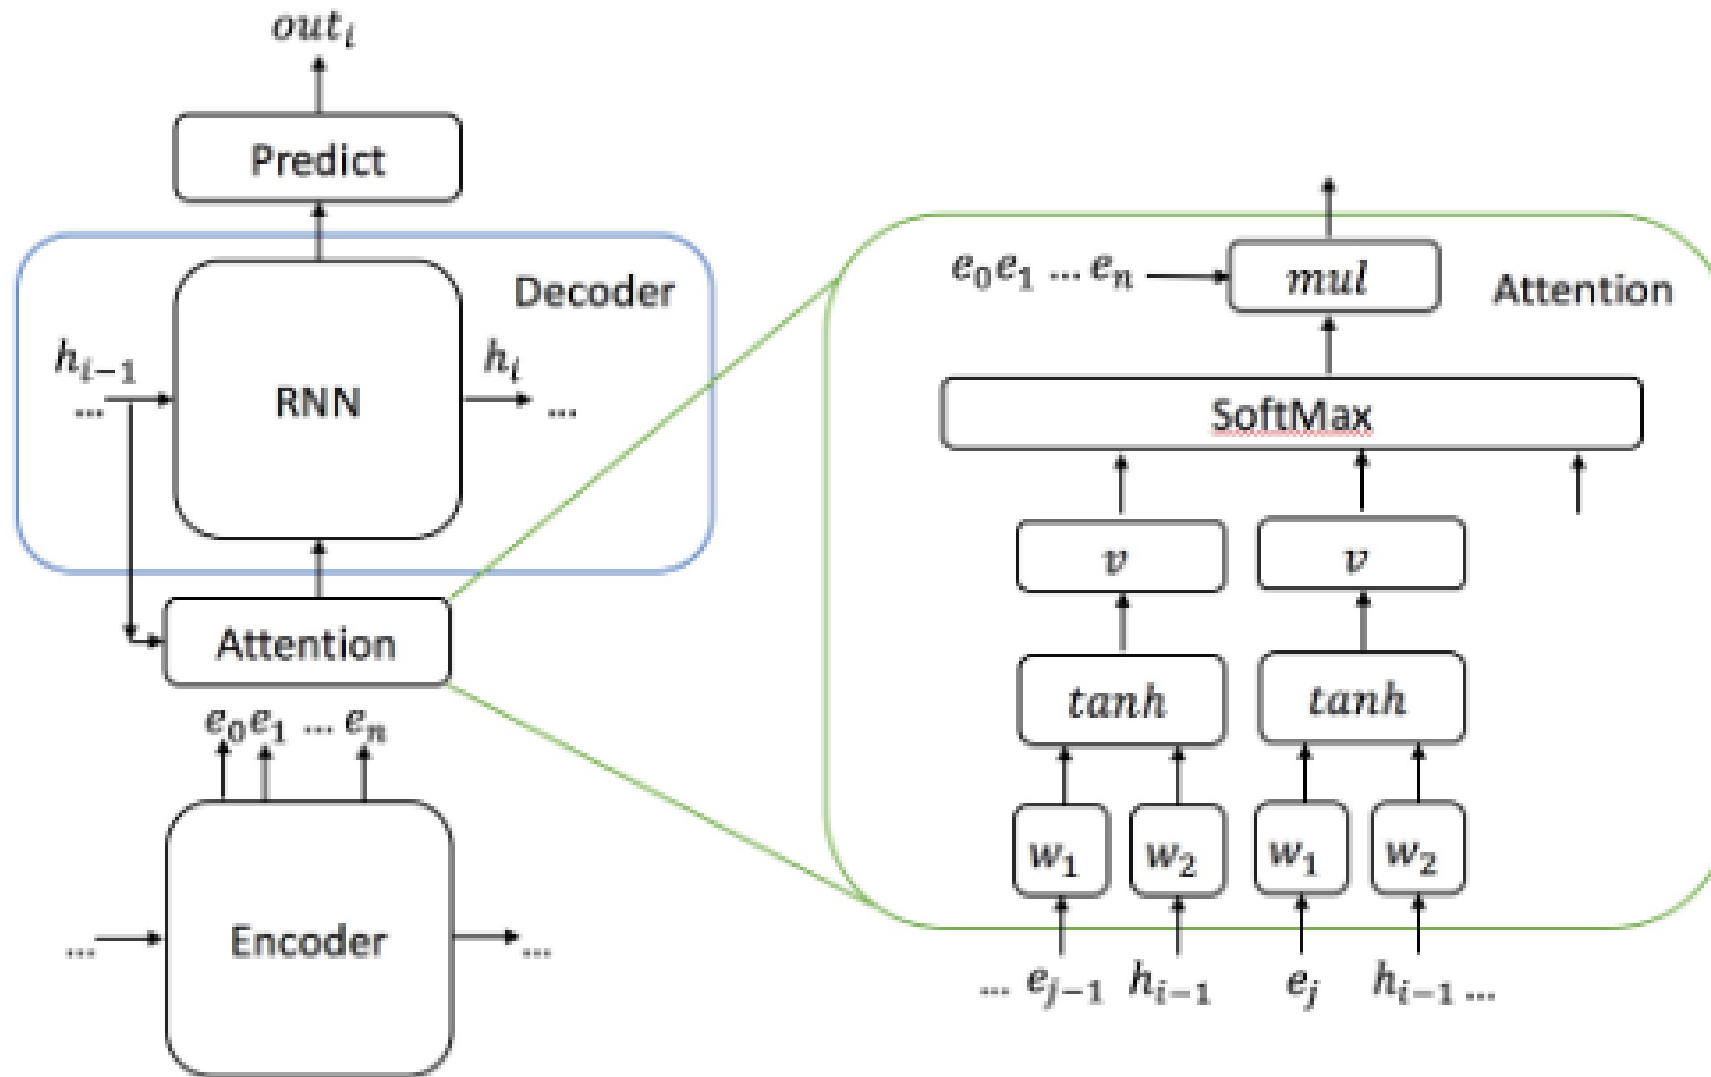


Decoder



Encoder

# Variations – Attention




# Papers

# Learning Phrase Representation using RNN Encoder-Decoder for Statistical Machine Translation



Kyunghyun Cho



Dzmitry Bahdanau



Fethi Bougares



Yoshua Bengio



Holger Schwenk

# Introduction

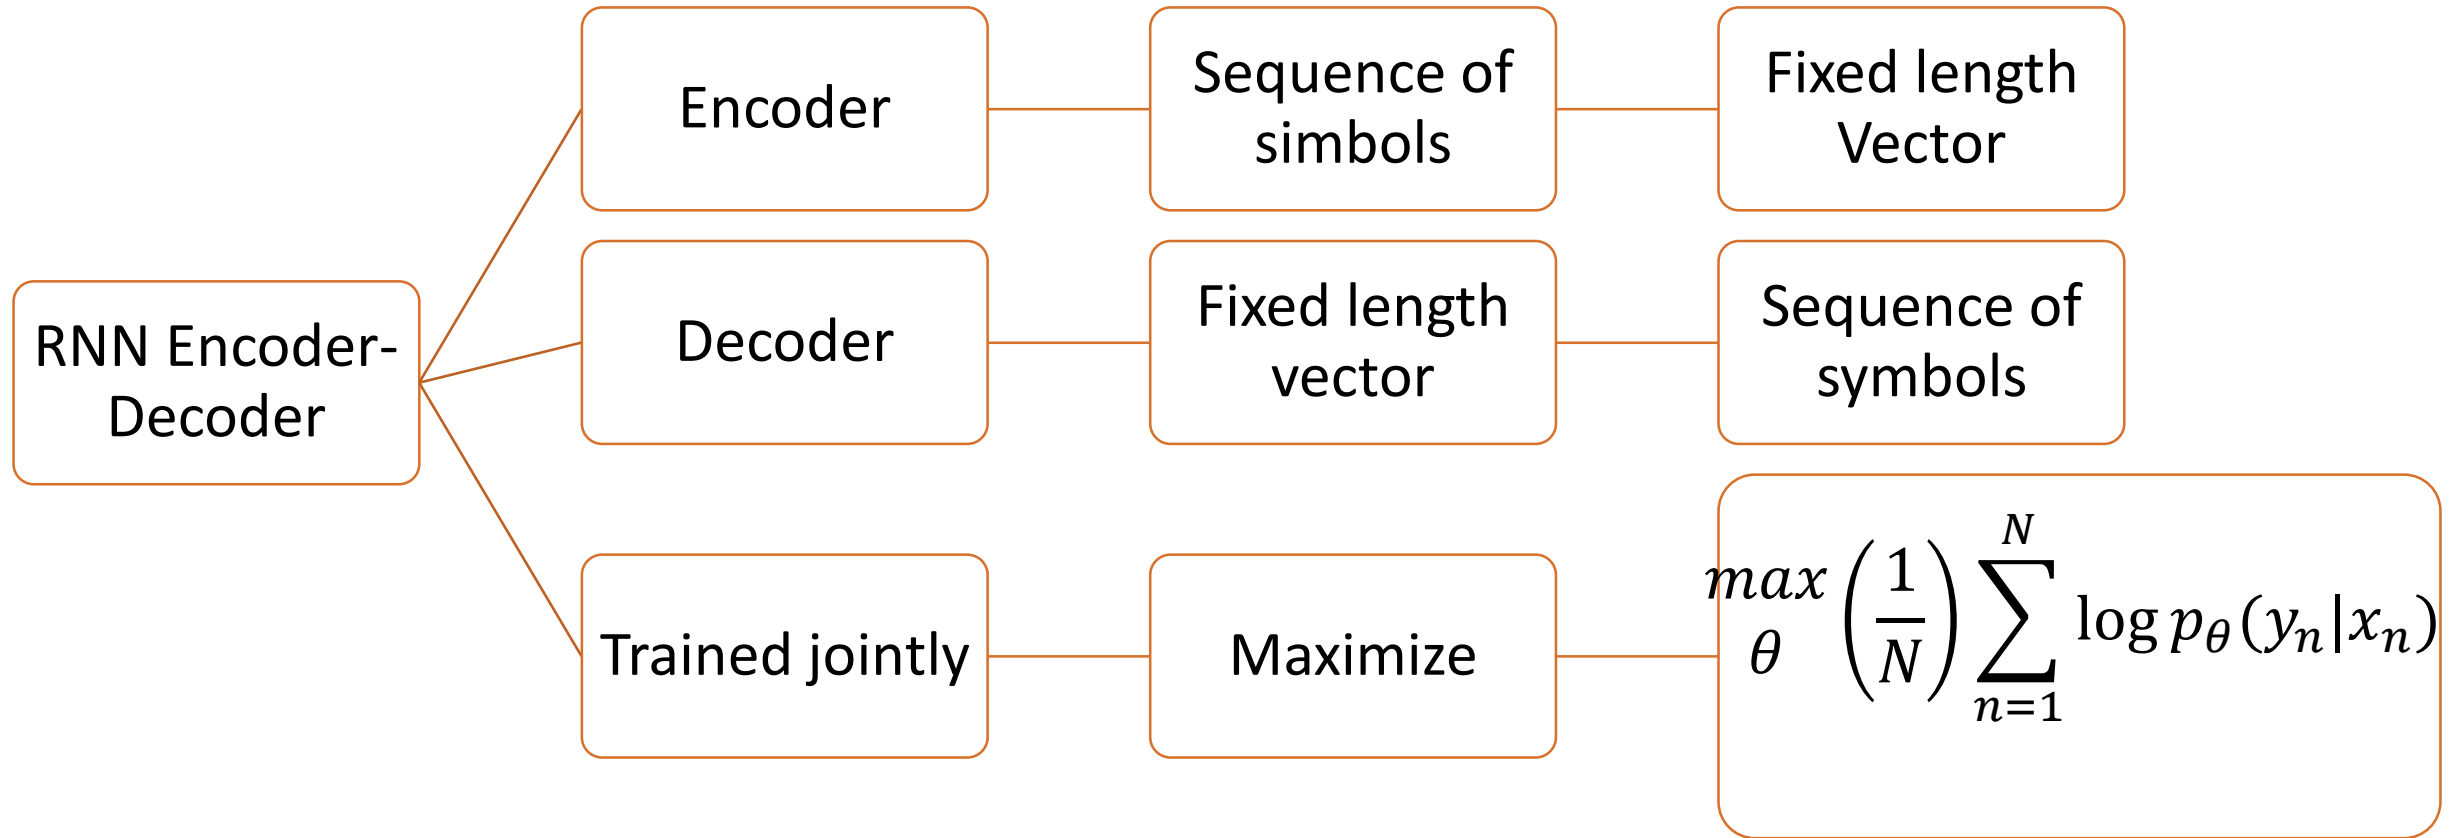
- NN used in natural language processing (NLP)
- Encoder – Decoder RNN for statistical Machine Translation (SMT)
- Sophisticated hidden unit
- Evaluation Context: translating English to French



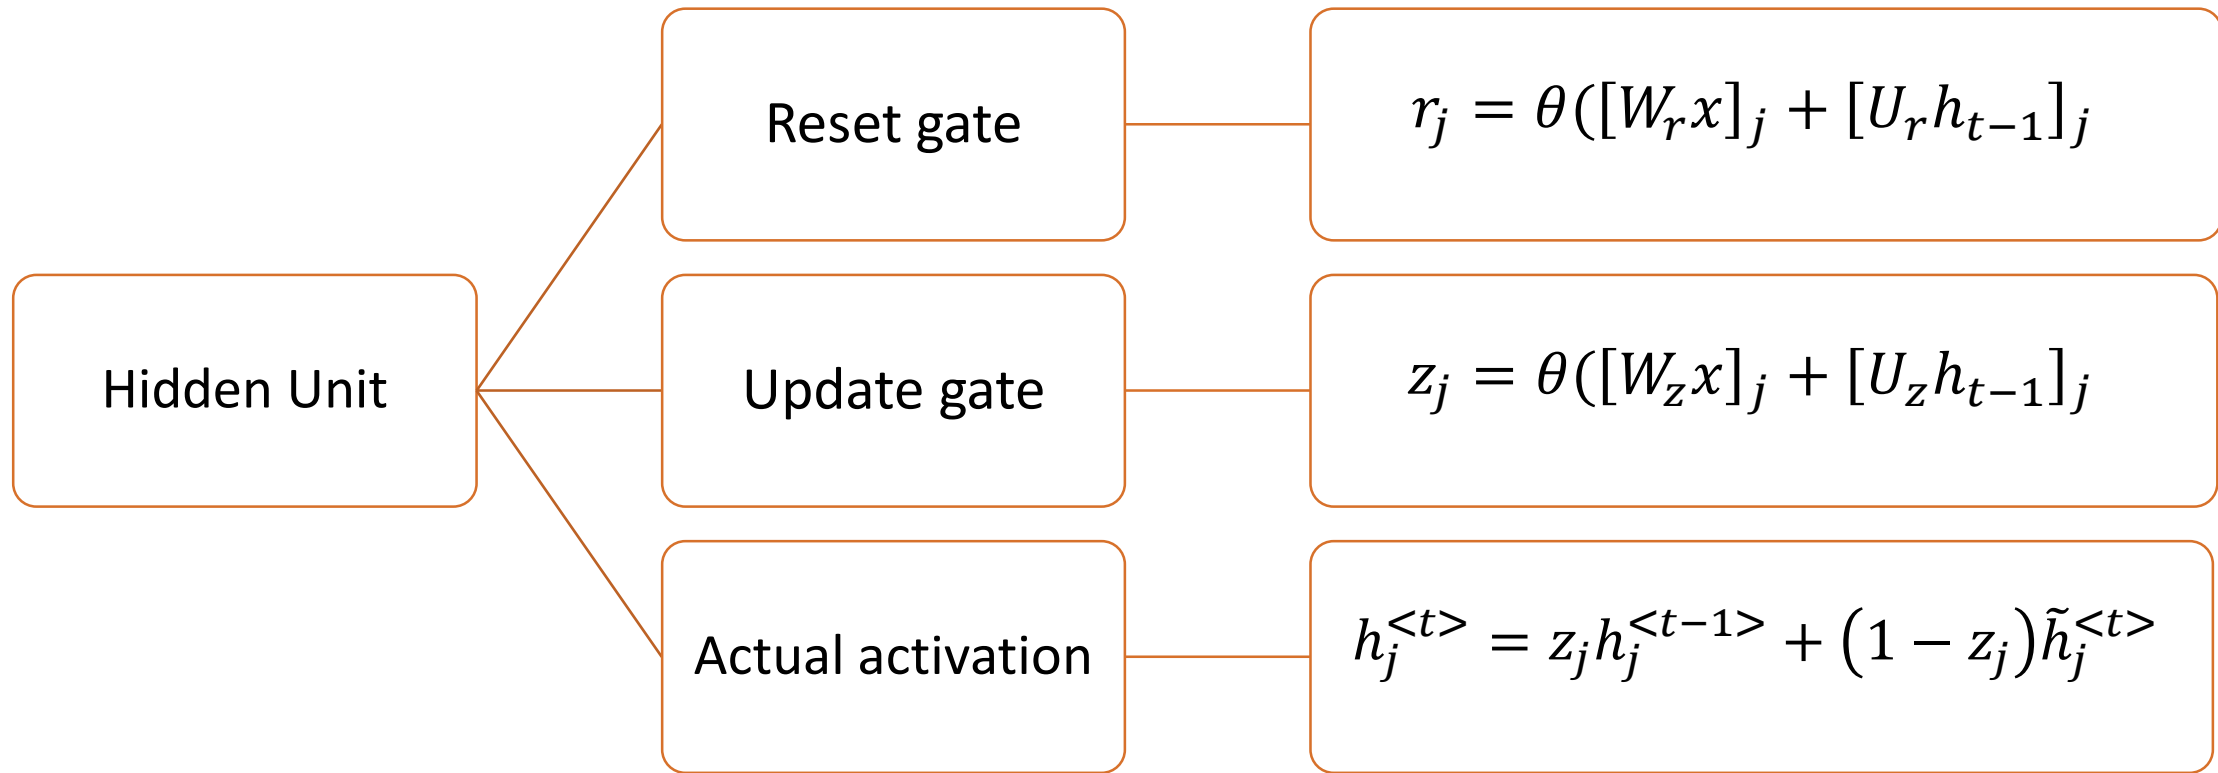
# Main Contributions

1. Encoder-Decoder structure using two RNN
2. GRU
3. Continuous Space Representation

# Proposed Methods

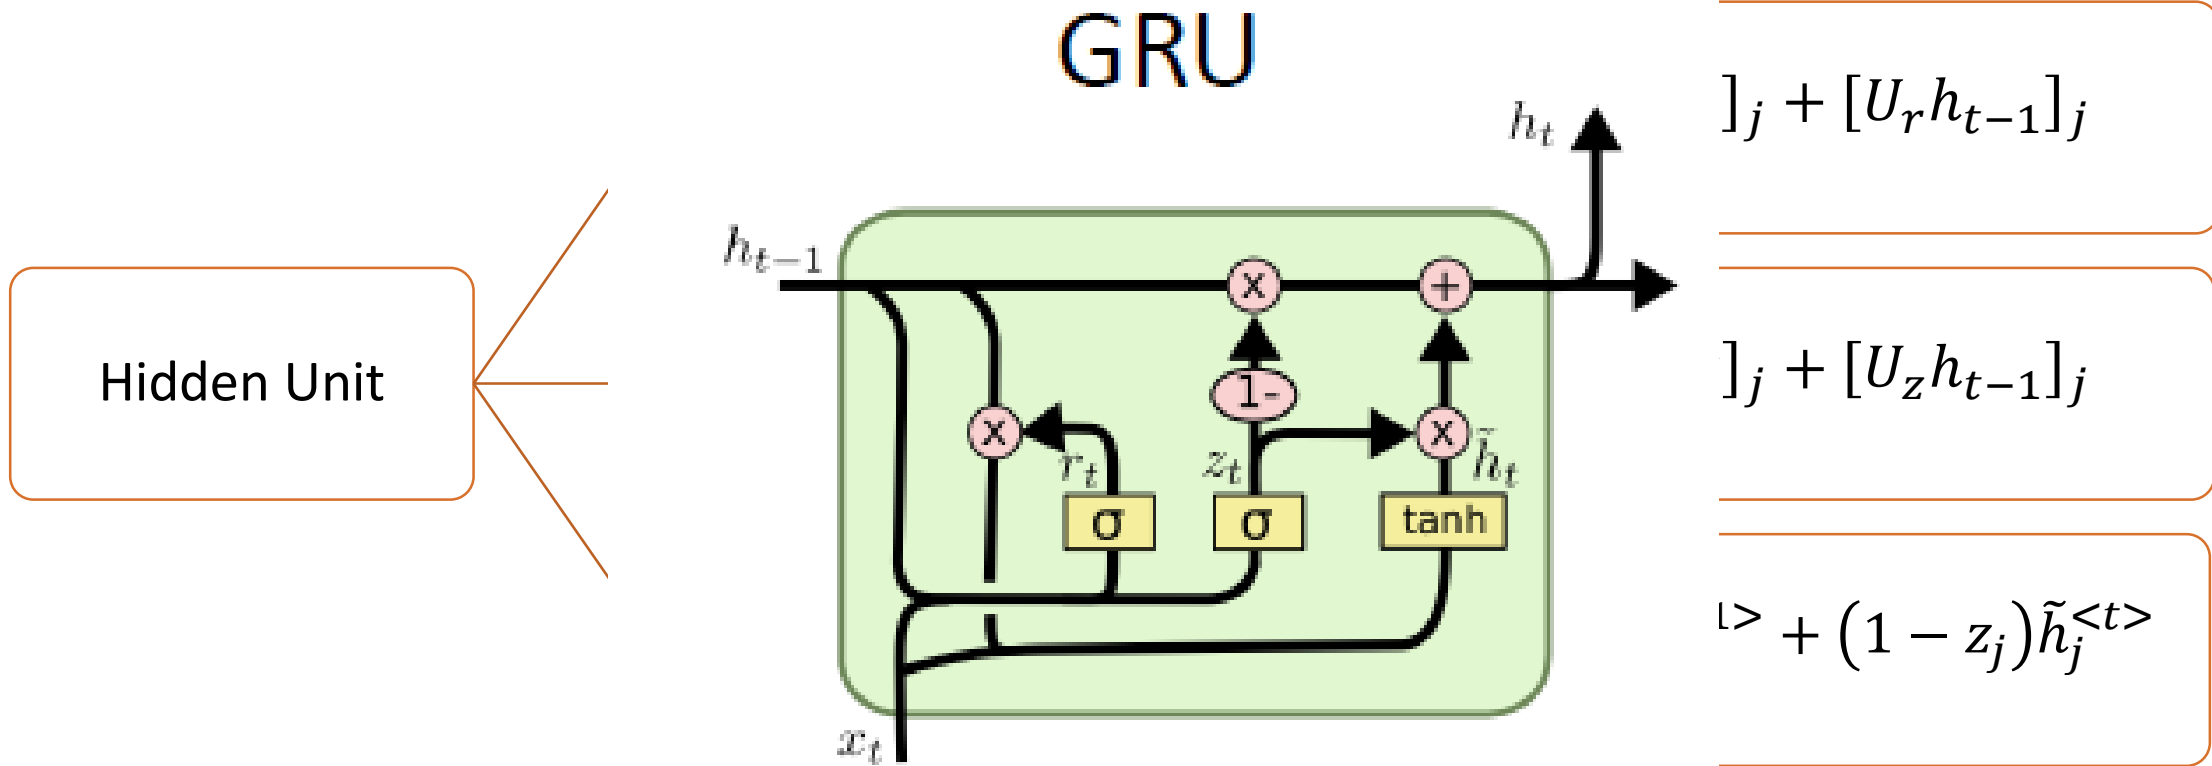


# Proposed Methods



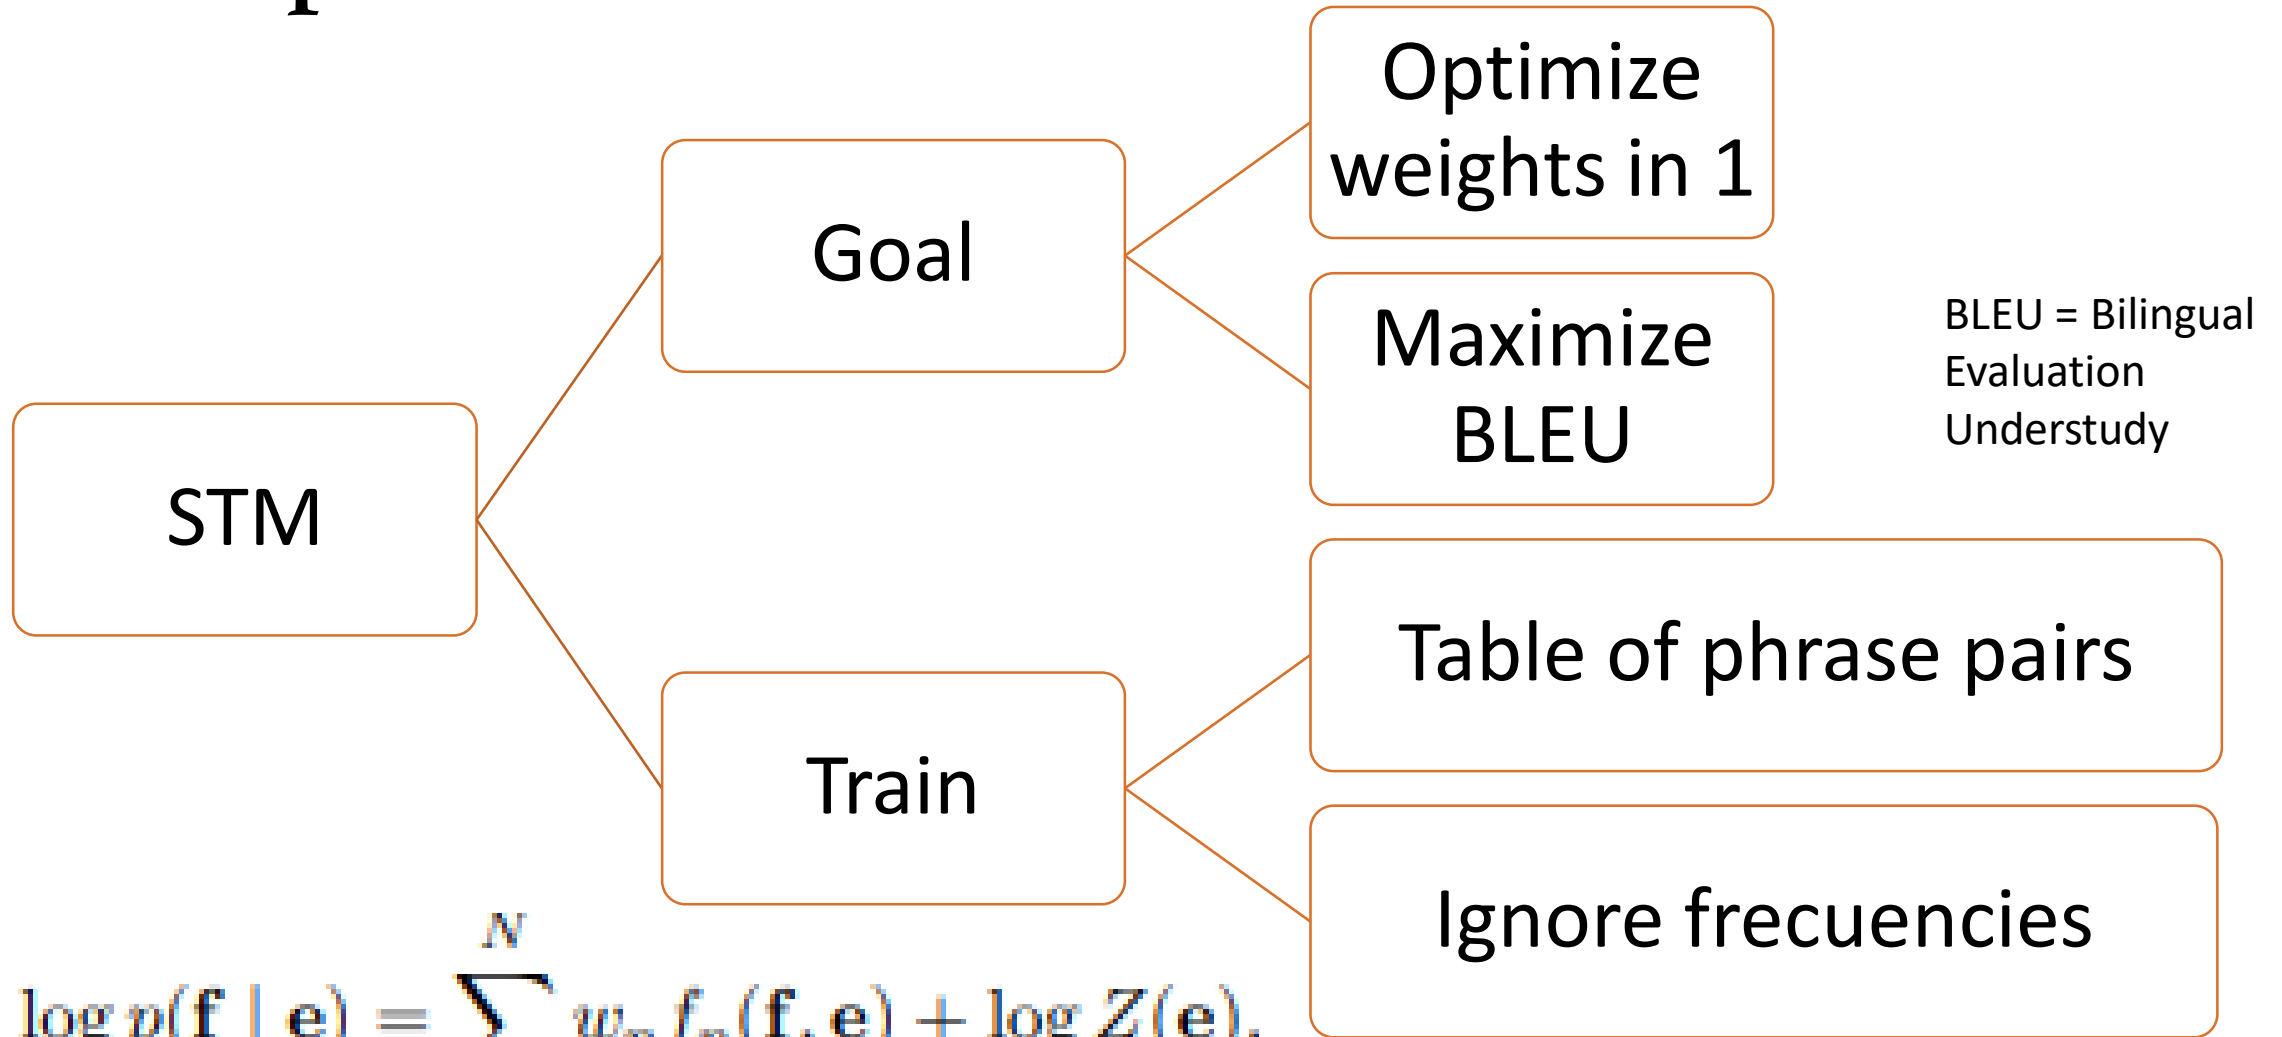
$$\tilde{h}_j^{<t>} = \phi \left( [\mathbf{W}\mathbf{x}]_j + [\mathbf{U} (\mathbf{r} \odot \mathbf{h}_{<t-1>})]_j \right)$$

# Proposed Methods



$$\tilde{h}_j^{<t>} = \phi \left( [Wx]_j + [U(r \odot h_{<t-1>})]_j \right)$$

# Proposed Methods



1.  $\log p(\mathbf{f} \mid \mathbf{e}) = \sum_{n=1}^N w_n f_n(\mathbf{f}, \mathbf{e}) + \log Z(\mathbf{e}),$

# Dataset

[ACL 2014 NINTH WORKSHOP  
ON STATISTICAL MACHINE TRANSLATION](#)

## **Shared Task: Machine Translation**

26-27 June 2014  
Baltimore, USA

[\[HOME\]](#) | [\[TRANSLATION TASK\]](#) | [\[METRICS TASK\]](#) | [\[QUALITY ESTIMATION TASK\]](#) |  
[\[MEDICAL TRANSLATION TASK\]](#) | [\[SCHEDULE\]](#) | [\[PAPERS\]](#) | [\[AUTHORS\]](#) | [\[RESULTS\]](#)

The recurring translation task of the [WMT workshops](#) focuses mainly on European language pairs, but this year we have introduced English-Hindi as an experimental, low resource language pair. Translation quality will be evaluated on a shared, unseen test set of news stories. We provide a parallel corpus as training data, a baseline system, and additional resources [for download](#). Participants may augment the baseline system or use their own system.

## **WMT' 14 workshop**

- Europarl
- News commentary
- UN
- Two crawled corpora
- Train set
  - Most frequent 15000 words
- Test set
  - Data selection (newstest2012 and 2013)
    - Weight tuning with MERT
  - newstest2014

# Time function representations of data

**TABLE 1.** Set of time functions considered in this work.

#	Feature
1	x-coordinate: $x_n$
2	y-coordinate: $y_n$
3	Pen-pressure: $z_n$
4	Path-tangent angle: $\theta_n$
5	Path velocity magnitude: $v_n$
6	Log curvature radius: $\rho_n$
7	Total acceleration magnitude: $a_n$
8-14	First-order derivate of features 1-7: $\dot{x}_n, \dot{y}_n, \dot{z}_n, \dot{\theta}_n, \dot{v}_n, \dot{\rho}_n, \dot{a}_n$
15-16	Second-order derivate of features 1-2: $\ddot{x}_n, \ddot{y}_n$
17	Ratio of the minimum over the maximum speed over a 5-samples window: $v_n^r$
18-19	Angle of consecutive samples and first order difference: $\alpha_n, \alpha'_n$
20	Sine: $s_n$
21	Cosine: $c_n$
22	Stroke length to width ratio over a 5-samples window: $r_n^5$
23	Stroke length to width ratio over a 7-samples window: $r_n^7$

# Quantitative Results

Models	BLEU	
	dev	test
Baseline	30.64	33.30
RNN	31.20	33.87
CSLM + RNN	31.48	34.64
CSLM + RNN + WP	31.50	34.54

Table 1: BLEU scores computed on the development and test sets using different combinations of approaches. WP denotes a *word penalty*, where we penalizes the number of unknown words to neural networks.



# Qualitative Results

Source	Translation Model	RNN Encoder-Decoder
at the end of the	[à la fin de la] [à la fin des années] [être supprimés à la fin de la]	[à la fin du] [à la fin des] [à la fin de la]
for the first time	[r © pour la première fois] [été donnés pour la première fois] [été commémorée pour la première fois]	[pour la première fois] [pour la première fois .] [pour la première fois que]
in the United States and	[? aux ?tats-Unis et] [été ouvertes aux États-Unis et] [été constatées aux États-Unis et]	[aux États-Unis et] [des États-Unis et] [des États-Unis et]
, as well as	[?s , qu"] [?s , ainsi que] [?re aussi bien que]	[, ainsi qu"] [, ainsi que] [, ainsi que les]
one of the most	[?t ?l' un des plus] [?l' un des plus] [être retenue comme un de ses plus]	[l' un des] [le] [un des]

(a) Long, frequent source phrases

Source	Translation Model	RNN Encoder-Decoder
, Minister of Communications and Transport	[Secrétaire aux communications et aux transports :] [Secrétaire aux communications et aux transports]	[Secrétaire aux communications et aux transports] [Secrétaire aux communications et aux transports :]
did not comply with the	[vestimentaire , ne correspondaient pas à des] [susmentionnée n' était pas conforme aux] [présentées n' étaient pas conformes à la]	[n' ont pas respecté les] [n' était pas conforme aux] [n' ont pas respecté la]
parts of the world .	[© gions du monde .] [régions du monde considérées .] [région du monde considérée .]	[parties du monde .] [les parties du monde .] [des parties du monde .]
the past few days .	[le petit texte .] [cours des tout derniers jours .] [les tout derniers jours .]	[ces derniers jours .] [les derniers jours .] [cours des derniers jours .]
on Friday and Saturday	[vendredi et samedi à la] [vendredi et samedi à] [se déroulera vendredi et samedi ,]	[le vendredi et le samedi] [le vendredi et samedi] [vendredi et samedi]

(b) Long, rare source phrases

Table 2: The top scoring target phrases for a small set of source phrases according to the translation model (direct translation probability) and by the RNN Encoder-Decoder. Source phrases were randomly selected from phrases with 4 or more words. ? denotes an incomplete (partial) character. r is a Cyrillic letter *ghe*.

# Qualitative Results

Source	Samples from RNN Encoder-Decoder
at the end of the	[à la fin de la] ( $\times 11$ )
for the first time	[pour la première fois] ( $\times 24$ ) [pour la première fois que] ( $\times 2$ )
in the United States and	[aux États-Unis et] ( $\times 6$ ) [dans les États-Unis et] ( $\times 4$ )
, as well as	[, ainsi que] [,] [ainsi que] [, ainsi qu'] [et UNK]
one of the most	[l' un des plus] ( $\times 9$ ) [l' un des] ( $\times 5$ ) [l' une des plus] ( $\times 2$ )

(a) Long, frequent source phrases

Source	Samples from RNN Encoder-Decoder
, Minister of Communica- tions and Transport	[, ministre des communications et le transport] ( $\times 13$ )
did not comply with the	[n' tait pas conforme aux] [n' a pas respect l'] ( $\times 2$ ) [n' a pas respect la] ( $\times 3$ )
parts of the world .	[arts du monde .] ( $\times 11$ ) [des arts du monde .] ( $\times 7$ )
the past few days .	[quelques jours .] ( $\times 5$ ) [les derniers jours .] ( $\times 5$ ) [ces derniers jours .] ( $\times 2$ )
on Friday and Saturday	[vendredi et samedi] ( $\times 5$ ) [le vendredi et samedi] ( $\times 7$ ) [le vendredi et le samedi] ( $\times 4$ )

(b) Long, rare source phrases

Table 3: Samples generated from the RNN Encoder-Decoder for each source phrase used in Table 2. We show the top-5 target phrases out of 50 samples. They are sorted by the RNN Encoder-Decoder scores.

# Word and Phrase Representation





Figure 4: 2-D embedding of the learned word representation. The left one shows the full embedding space, while the right one shows a zoomed-in view of one region (color-coded). For more plots, see the supplementary material.

# Word and Phrase Representation

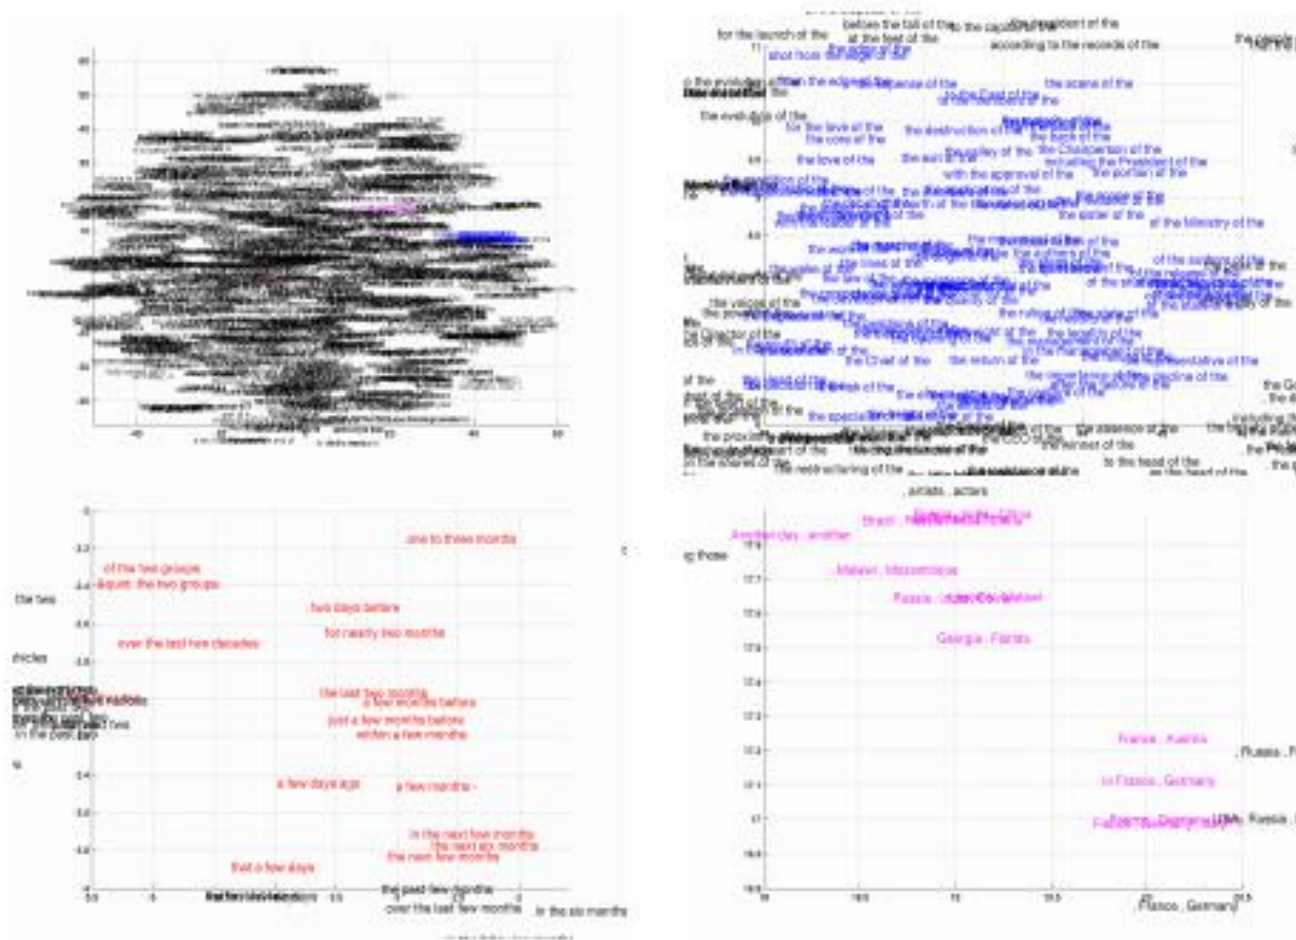


Figure 5: 2-D embedding of the learned phrase representation. The top left one shows the full representation space (5000 randomly selected points), while the other three figures show the zoomed-in view of specific regions (color-coded).

# 4 vs. 1 – Average pairs scores


**TABLE 3. 4vs1 Evaluation Results:** System performance in terms of EER(%) for the three different training scenarios considered, i.e., “skilled”, “random” and “skilled + random”.

	Train: “skilled”		Train: “random”		Train: “skilled + random”	
	Skilled	Random	Skilled	Random	Skilled	Random
LSTM	5.58	24.03	15.17	4.08	6.17	3.67
GRU	6.25	28.69	13.92	4.25	5.58	3.63
BLSTM	<b>4.75</b>	24.03	15.58	3.89	<b>5.50</b>	3.00
BGRU	4.92	19.69	12.33	<b>3.25</b>	5.92	<b>2.92</b>

**TABLE 4. 1vs1 and 4vs1 DTW-based Evaluation Results:** System performance in terms of EER(%).

	1vs1	4vs1
Skilled	10.17	7.75
Random	0.94	0.50

# Detection Error Tradeoff curve



**FIGURE 7.** System performance results obtained using our Proposed BLSTM System for the 4vs1 case and “skilled + random” train scenario over the BiosecuID evaluation dataset.

To achieve a state-of-the-art performance of the model for both skilled and random forgeries, a possible solution is to perform two consecutive stages:

1. Stage based on DTW optimized for rejecting random forgeries.
2. Proposed RNNs Systems in order to reject the remaining skilled forgeries.

# Conclusions

- RNN Encoder-Decoder is able to learn mapping from a sequence to another.
- Also, is able to score a pair of sequences or generate a target given a source sequence.
- The hidden unit is able to adaptively control how much it remembers or forget while reading or generating a sequence
- The model is able to capture linguistic regularities.
- The RNN Encoder–Decoder is able to propose well-formed target phrases
- The RNN Encoder–Decoder improves BLEU.
- Potential for improvement and analysis!

# Exploring Recurrent Neural Networks for On-Line Handwritten Signature Biometrics



Ruben tolosana



Ruben vera-rodriquez



Julian fierrez



Javier ortega-garcia



# Introduction

- Bidirectional LSTMs and GRUs RNNs caused great impact in handwriting recognition due to the relationship that exists between current inputs and past and future contexts.
- Off-line vs. On-line considerations
- Despite the good results obtained in the field of handwriting recognition, very few studies have successfully RNN architectures to handwritten signature verification.


# Introduction

- Until now LSTM RNN systems trained with standard mechanisms are not appropriate for the task of signature verification as the amount of available data for this task is scarce.

# Main Contributions

1. RNNs with a Siamese architecture
2. Writer-independent scenario
3. Strict experimental protocol
4. First analysis of RNNs for the two types of forgeries considered in on-line signature verification (i.e. skilled and random or zero-effort forgeries).
5. Bidirectional Scheme

# Proposed Methods



- Siamese Architecture
- LSTMs
- GRUs
- Bidirectional RNNs

**FIGURE 1.** Examples of our proposed LSTM and GRU RNN systems based on a Siamese architecture for minimizing a discriminative cost function. (a) Genuine case. (b) Impostor case.

# Dataset



## BiosecurID


- 400 Users
  - 16 original signature
  - 12 skilled forgeries signatures
  - 4 acquisition sessions
- Data of each signature
  - X and Y coordinates: 0.25 mm resolution
  - Pressure: 1024 levels
  - Timestamp: 100Hz

# Time function representations of data

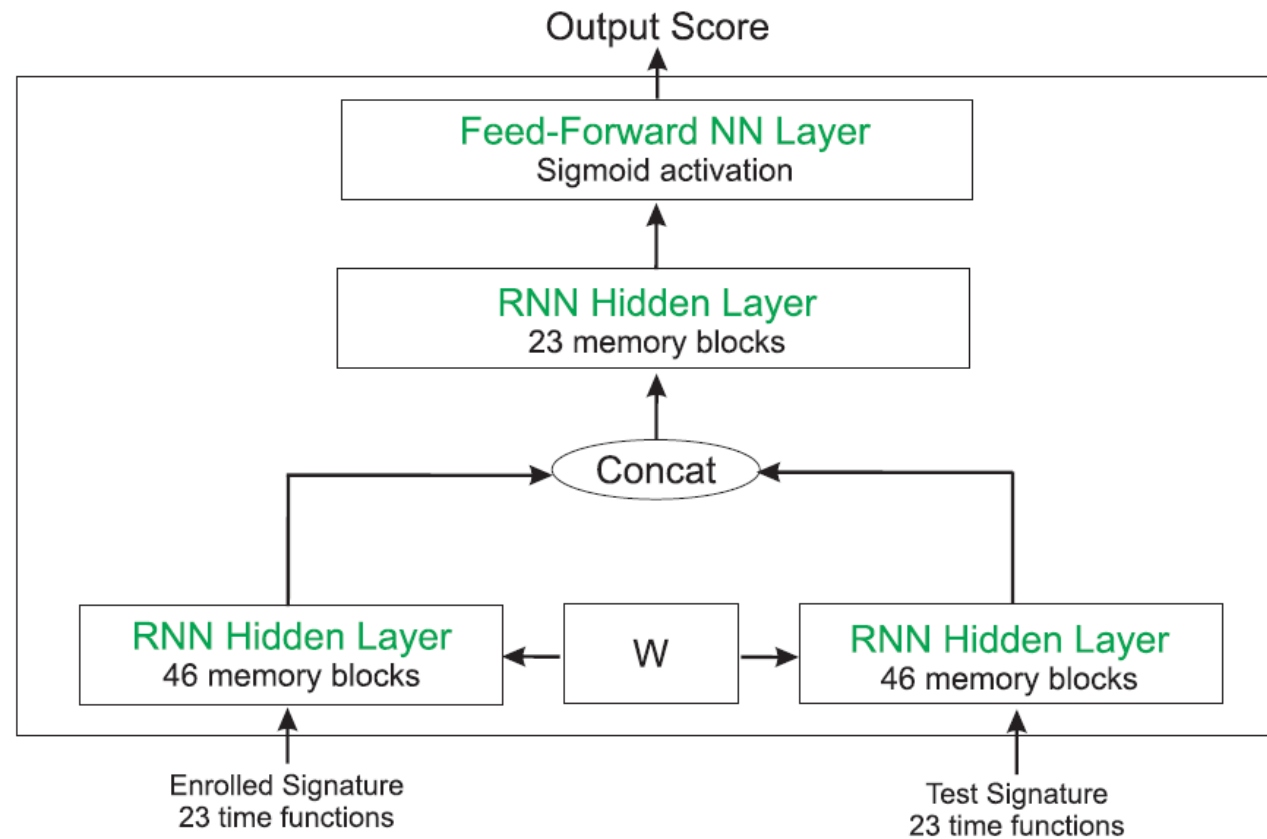
**TABLE 1.** Set of time functions considered in this work.

#	Feature
1	x-coordinate: $x_n$
2	y-coordinate: $y_n$
3	Pen-pressure: $z_n$
4	Path-tangent angle: $\theta_n$
5	Path velocity magnitude: $v_n$
6	Log curvature radius: $\rho_n$
7	Total acceleration magnitude: $a_n$
8-14	First-order derivate of features 1-7: $\dot{x}_n, \dot{y}_n, \dot{z}_n, \dot{\theta}_n, \dot{v}_n, \dot{\rho}_n, \dot{a}_n$
15-16	Second-order derivate of features 1-2: $\ddot{x}_n, \ddot{y}_n$
17	Ratio of the minimum over the maximum speed over a 5-samples window: $v_n^r$
18-19	Angle of consecutive samples and first order difference: $\alpha_n, \alpha'_n$
20	Sine: $s_n$
21	Cosine: $c_n$
22	Stroke length to width ratio over a 5-samples window: $r_n^5$
23	Stroke length to width ratio over a 7-samples window: $r_n^7$

# Experimental Protocol




# Final topology



**FIGURE 5.** End-to-end on-line signature verification system proposed in this work and based on the use of LSTM and GRU RNNs with a Siamese architecture.



# Training Cost



**FIGURE 6.** Considered RNNs cost during training for the “skilled” scenario. A small green vertical line indicates for each proposed RNN system the training iteration which provides the best system performance over the evaluation dataset.

# 1 vs. 1 – All pair scores

**TABLE 2. 1vs1 Evaluation Results:** System performance in terms of EER(%) for the three different training scenarios considered, i.e., “skilled”, “random” and “skilled + random”.

	Train: “skilled”		Train: “random”		Train: “skilled + random”	
	Skilled	Random	Skilled	Random	Skilled	Random
LSTM	6.44	24.48	13.31	5.38	7.94	6.22
GRU	7.69	29.42	15.63	6.92	7.67	5.98
BLSTM	<b>5.60</b>	24.48	15.31	<b>5.28</b>	<b>6.83</b>	<b>5.38</b>
BGRU	6.31	19.14	12.56	5.33	7.88	5.52

**TABLE 4. 1vs1 and 4vs1 DTW-based Evaluation Results:** System performance in terms of EER(%).

	1vs1	4vs1
Skilled	10.17	7.75
Random	0.94	0.50

# 4 vs. 1 – Average pairs scores


**TABLE 3. 4vs1 Evaluation Results:** System performance in terms of EER(%) for the three different training scenarios considered, i.e., “skilled”, “random” and “skilled + random”.

	Train: “skilled”		Train: “random”		Train: “skilled + random”	
	Skilled	Random	Skilled	Random	Skilled	Random
LSTM	5.58	24.03	15.17	4.08	6.17	3.67
GRU	6.25	28.69	13.92	4.25	5.58	3.63
BLSTM	<b>4.75</b>	24.03	15.58	3.89	<b>5.50</b>	3.00
BGRU	4.92	19.69	12.33	<b>3.25</b>	5.92	<b>2.92</b>

**TABLE 4. 1vs1 and 4vs1 DTW-based Evaluation Results:** System performance in terms of EER(%).

	1vs1	4vs1
Skilled	10.17	7.75
Random	0.94	0.50

# DET curve



**FIGURE 7.** System performance results obtained using our Proposed BLSTM System for the 4vs1 case and “skilled + random” train scenario over the BiosecuID evaluation dataset.

To achieve a state-of-the-art performance of the model for both skilled and random forgeries, a possible solution is to perform two consecutive stages:

1. Stage based on DTW optimized for rejecting random forgeries.
2. Proposed RNNs Systems in order to reject the remaining skilled forgeries.

# Conclusions

- The main contribution is to assess the feasibility of different RNNs systems in combination with a Siamese architecture for on-line handwritten signature verification.
- First complete and successful framework on the use of multiple RNN systems (i.e. LSTM and GRU) for on-line handwritten signature verification considering both skilled and random types of forgeries.
- Difference in the number of training iterations needed between normal and bidirectional schemes.
- Difference in the number of training iterations between both LSTM and GRU RNNs.
- High ability of our proposed approach for learning even with small amounts of signatures.

# Tutorial

```
from __future__ import unicode_literals, print_function, division
from io import open
import unicodedata
import string
import re
import random

import torch
import torch.nn as nn
from torch import optim
import torch.nn.functional as F

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
```

```
SOS_token = 0
EOS_token = 1
```

```
class Lang:
    def __init__(self, name):
        self.name = name
        self.word2index = {}
        self.word2count = {}
        self.index2word = {0: "SOS", 1: "EOS"}
        self.n_words = 2 # Count SOS and EOS

    def addSentence(self, sentence):
        for word in sentence.split(' '):
            self.addWord(word)

    def addWord(self, word):
        if word not in self.word2index:
            self.word2index[word] = self.n_words
            self.word2count[word] = 1
            self.index2word[self.n_words] = word
            self.n_words += 1
        else:
            self.word2count[word] += 1
```

1. We will be representing each word in a language as a one-hot vector.
2. We will however use a few thousand words per language.
3. We'll need a unique index per word to use as the inputs and targets of the networks later.
4. To keep track of all this we will use a helper class called Lang



```
# Turn a Unicode string to plain ASCII, thanks to  
# https://stackoverflow.com/a/518232/2809427  
def unicodeToAscii(s):  
    return ''.join(  
        c for c in unicodedata.normalize('NFD', s)  
        if unicodedata.category(c) != 'Mn'  
    )  
  
# Lowercase, trim, and remove non-letter characters  
  
def normalizeString(s):  
    s = unicodeToAscii(s.lower().strip())  
    s = re.sub(r"([.!?])", r" \1", s)  
    s = re.sub(r"[^a-zA-Z.!?]+", r" ", s)  
    return s
```

1. Convert Unicode to ASCII.
2. Make all words Lowercase, trim non-letter characters.

```

def readLangs(lang1, lang2, reverse=False):
    print("Reading lines...")

    # Read the file and split into lines
    lines = open('data/%s-%s.txt' % (lang1, lang2), encoding='utf-8').\
        read().strip().split('\n')

    # Split every line into pairs and normalize
    pairs = [[normalizeString(s) for s in l.split('\t')] for l in lines]

    # Reverse pairs, make Lang instances
    if reverse:
        pairs = [list(reversed(p)) for p in pairs]
        input_lang = Lang(lang2)
        output_lang = Lang(lang1)
    else:
        input_lang = Lang(lang1)
        output_lang = Lang(lang2)

    return input_lang, output_lang, pairs

```

1. Read the file.
2. Split it in pairs [ENG-FRE].

```
MAX_LENGTH = 10
```

```
eng_prefixes = (  
    "i am ", "i m ",  
    "he is", "he s ",  
    "she is", "she s ",  
    "you are", "you re ",  
    "we are", "we re ",  
    "they are", "they re "  
)
```

```
def filterPair(p):  
    return len(p[0].split(' ')) < MAX_LENGTH and \  
        len(p[1].split(' ')) < MAX_LENGTH and \  
        p[1].startswith(eng_prefixes)
```

```
def filterPairs(pairs):  
    return [pair for pair in pairs if filterPair(pair)]
```

1. Filter phrases by size.
2. Filter phrases by prefixes.

```
def prepareData(lang1, lang2, reverse=False):
    input_lang, output_lang, pairs = readLangs(lang1, lang2, reverse)
    print("Read %s sentence pairs" % len(pairs))
    pairs = filterPairs(pairs)
    print("Trimmed to %s sentence pairs" % len(pairs))
    print("Counting words...")
    for pair in pairs:
        input_lang.addSentence(pair[0])
        output_lang.addSentence(pair[1])
    print("Counted words:")
    print(input_lang.name, input_lang.n_words)
    print(output_lang.name, output_lang.n_words)
    return input_lang, output_lang, pairs

input_lang, output_lang, pairs = prepareData('eng', 'fra', True)
print(random.choice(pairs))
```

1. Prepare all the data using the previous functions.

```

class EncoderRNN(nn.Module):
    def __init__(self, input_size, hidden_size):
        super(EncoderRNN, self).__init__()
        self.hidden_size = hidden_size

        self.embedding = nn.Embedding(input_size, hidden_size)
        self.gru = nn.GRU(hidden_size, hidden_size)

    def forward(self, input, hidden):
        embedded = self.embedding(input).view(1, 1, -1)
        output = embedded
        output, hidden = self.gru(output, hidden)
        return output, hidden

    def initHidden(self):
        return torch.zeros(1, 1, self.hidden_size, device=device)

```

1. Definition of the Encoder.
  - From a phrase to a vector representing the meaning of it.

```

class DecoderRNN(nn.Module):
    def __init__(self, hidden_size, output_size):
        super(DecoderRNN, self).__init__()
        self.hidden_size = hidden_size

        self.embedding = nn.Embedding(output_size, hidden_size)
        self.gru = nn.GRU(hidden_size, hidden_size)
        self.out = nn.Linear(hidden_size, output_size)
        self.softmax = nn.LogSoftmax(dim=1)

    def forward(self, input, hidden):
        output = self.embedding(input).view(1, 1, -1)
        output = F.relu(output)
        output, hidden = self.gru(output, hidden)
        output = self.softmax(self.out(output[0]))
        return output, hidden

    def initHidden(self):
        return torch.zeros(1, 1, self.hidden_size, device=device)

```

1. Definition of the Decoder.  
- From Hidden State of the Encoder to the translation.

```

class AttnDecoderRNN(nn.Module):
    def __init__(self, hidden_size, output_size, dropout_p=0.1, max_length=MAX_LENGTH):
        super(AttnDecoderRNN, self).__init__()
        self.hidden_size = hidden_size
        self.output_size = output_size
        self.dropout_p = dropout_p
        self.max_length = max_length

        self.embedding = nn.Embedding(self.output_size, self.hidden_size)
        self.attn = nn.Linear(self.hidden_size * 2, self.max_length)
        self.attn_combine = nn.Linear(self.hidden_size * 2, self.hidden_size)
        self.dropout = nn.Dropout(self.dropout_p)
        self.gru = nn.GRU(self.hidden_size, self.hidden_size)
        self.out = nn.Linear(self.hidden_size, self.output_size)

    def forward(self, input, hidden, encoder_outputs):
        embedded = self.embedding(input).view(1, 1, -1)
        embedded = self.dropout(embedded)

        attn_weights = F.softmax(
            self.attn(torch.cat((embedded[0], hidden[0]), 1))), dim=1)
        attn_applied = torch.bmm(attn_weights.unsqueeze(0),
                                encoder_outputs.unsqueeze(0))

        output = torch.cat((embedded[0], attn_applied[0]), 1)
        output = self.attn_combine(output).unsqueeze(0)

        output = F.relu(output)
        output, hidden = self.gru(output, hidden)

        output = F.log_softmax(self.out(output[0]), dim=1)
        return output, hidden, attn_weights

    def initHidden(self):
        return torch.zeros(1, 1, self.hidden_size, device=device)

```

1. Definition of the an Attention Decoder.
  - “Focus” on a different part of the encoder’s outputs for every step of the decoder’s own outputs.

```

def indexesFromSentence(lang, sentence):
    return [lang.word2index[word] for word in sentence.split(' ')]

def tensorFromSentence(lang, sentence):
    indexes = indexesFromSentence(lang, sentence)
    indexes.append(EOS_token)
    return torch.tensor(indexes, dtype=torch.long, device=device).view(-1, 1)

def tensorsFromPair(pair):
    input_tensor = tensorFromSentence(input_lang, pair[0])
    target_tensor = tensorFromSentence(output_lang, pair[1])
    return (input_tensor, target_tensor)

```

1. Get the indices that represents a sentence.
2. Get a tensor that represent the sentence.
3. Get a tensors that represent the input phrase and another that represents the target phrase.



```

teacher_forcing_ratio = 0.5

def train(input_tensor, target_tensor, encoder, decoder, encoder_optimizer, decoder_optimizer,
criterion, max_length=MAX_LENGTH):
    encoder_hidden = encoder.initHidden()

    encoder_optimizer.zero_grad()
    decoder_optimizer.zero_grad()

    input_length = input_tensor.size(0)
    target_length = target_tensor.size(0)

    encoder_outputs = torch.zeros(max_length, encoder.hidden_size, device=device)

    loss = 0

    for ei in range(input_length):
        encoder_output, encoder_hidden = encoder(
            input_tensor[ei], encoder_hidden)
        encoder_outputs[ei] = encoder_output[0, 0]

    decoder_input = torch.tensor([[SOS_token]], device=device)

    decoder_hidden = encoder_hidden

    use_teacher_forcing = True if random.random() < teacher_forcing_ratio else False

    if use_teacher_forcing:
        # Teacher forcing: Feed the target as the next input
        for di in range(target_length):
            decoder_output, decoder_hidden, decoder_attention = decoder(
                decoder_input, decoder_hidden, encoder_outputs)
            loss += criterion(decoder_output, target_tensor[di])
            decoder_input = target_tensor[di] # Teacher forcing

    else:
        # Without teacher forcing: use its own predictions as the next input
        for di in range(target_length):
            decoder_output, decoder_hidden, decoder_attention = decoder(
                decoder_input, decoder_hidden, encoder_outputs)
            topv, topi = decoder_output.topk(1)
            decoder_input = topi.squeeze().detach() # detach from history as input

            loss += criterion(decoder_output, target_tensor[di])
            if decoder_input.item() == EOS_token:
                break

    loss.backward()

    encoder_optimizer.step()
    decoder_optimizer.step()

    return loss.item() / target_length

```

## 1. Train for one pair of inputs.

```

def trainIters(encoder, decoder, n_iters, print_every=1000, plot_every=100,
learning_rate=0.01):
    start = time.time()
    plot_losses = []
    print_loss_total = 0 # Reset every print_every
    plot_loss_total = 0 # Reset every plot_every

    encoder_optimizer = optim.SGD(encoder.parameters(), lr=learning_rate)
    decoder_optimizer = optim.SGD(decoder.parameters(), lr=learning_rate)
    training_pairs = [tensorsFromPair(random.choice(pairs))
                      for i in range(n_iters)]
    criterion = nn.NLLLoss()

    for iter in range(1, n_iters + 1):
        training_pair = training_pairs[iter - 1]
        input_tensor = training_pair[0]
        target_tensor = training_pair[1]

        loss = train(input_tensor, target_tensor, encoder,
                     decoder, encoder_optimizer, decoder_optimizer, criterion)
        print_loss_total += loss
        plot_loss_total += loss

        if iter % print_every == 0:
            print_loss_avg = print_loss_total / print_every
            print_loss_total = 0
            print('%s (%d %d%%) %.4f' % (timeSince(start, iter / n_iters),
                                         iter, iter / n_iters * 100, print_loss_avg))

        if iter % plot_every == 0:
            plot_loss_avg = plot_loss_total / plot_every
            plot_losses.append(plot_loss_avg)
            plot_loss_total = 0

    showPlot(plot_losses)

```

## 1. Training the model.

```
hidden_size = 256
encoder1 = EncoderRNN(input_lang.n_words, hidden_size).to(device)
attn_decoder1 = AttnDecoderRNN(hidden_size, output_lang.n_words, dropout_p=0.1).to(device)

trainIters(encoder1, attn_decoder1, 75000, print_every=5000)

output_words, attentions = evaluate(
    encoder1, attn_decoder1, "je suis trop froid .")
```

1. Define the size of the hidden state.
2. Create the encoder.
3. Create the decoder.
4. Train a model.
5. Evaluate the model with a phrase

# Homework

- Implement LSTM and SRU instead of GRU. Compare and discuss the results.
- Modify the architecture adding more layers. Show your results and discuss.
- Modify the RNN direction into a bidirectional network. Explain how this affect your initial results.